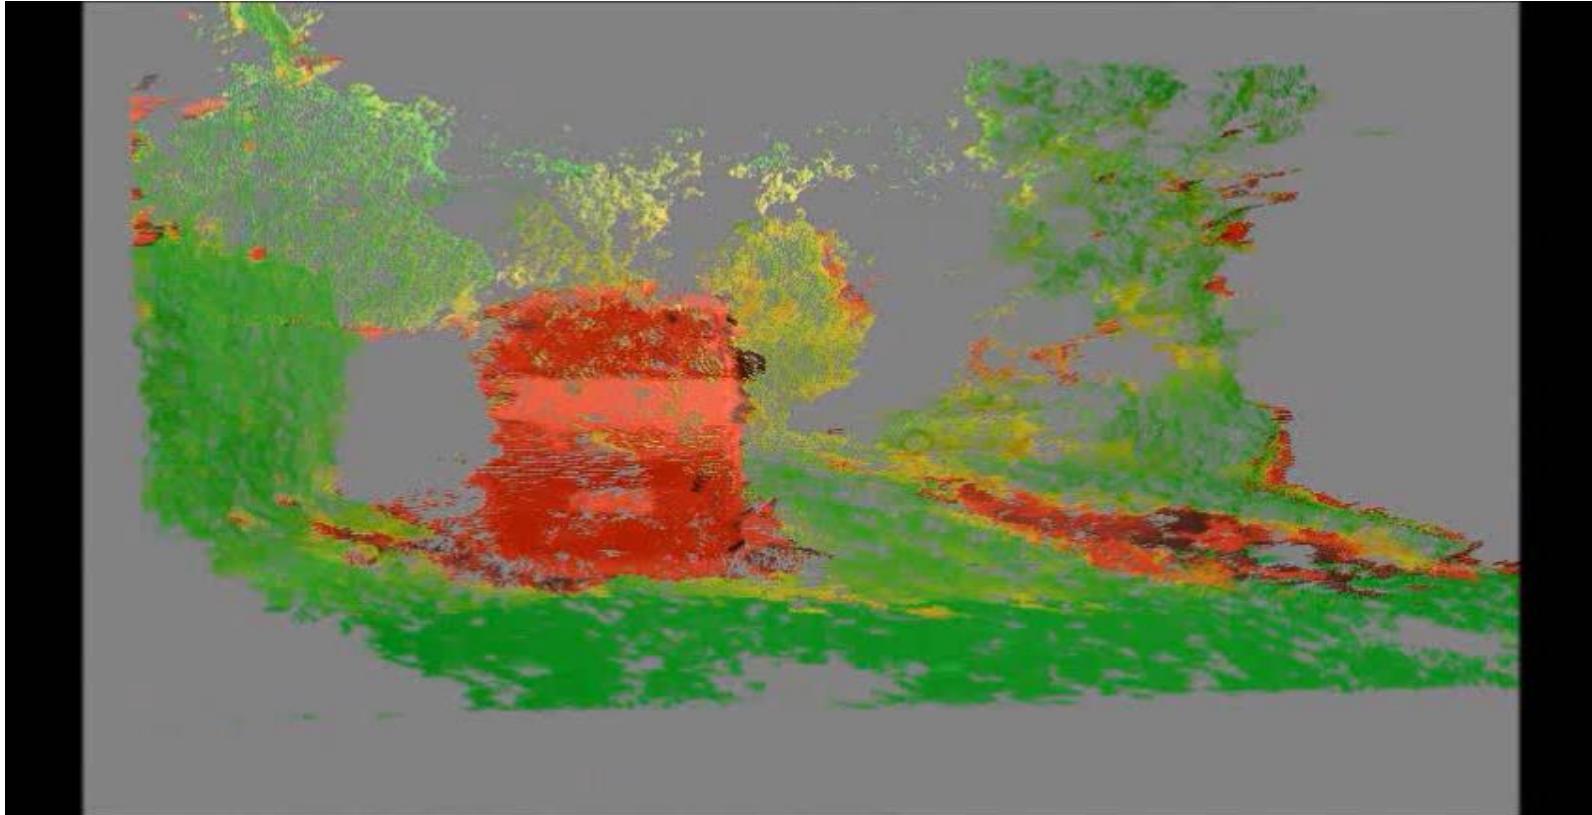


Introduction to Deep Learning

AI for Driver Assistance



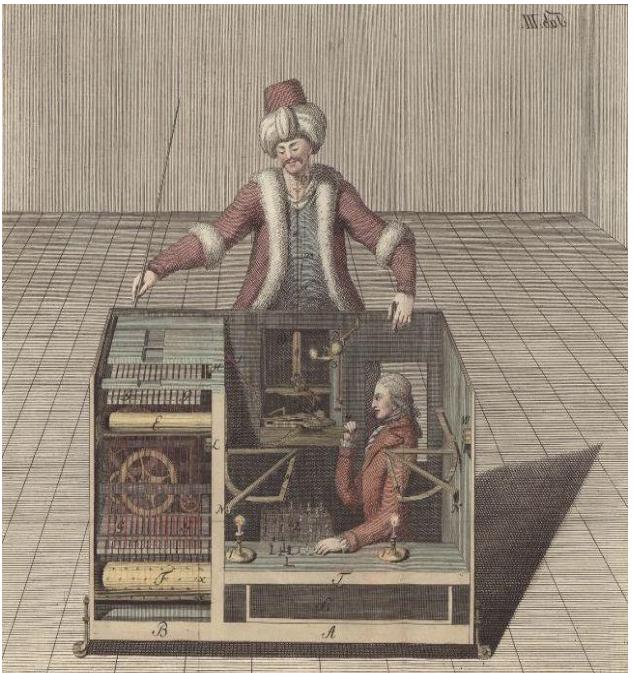
AI for Driver Assistance



Wedel et al. ECCV '08, Wedel, Cremers, Springer 2011

Introduction to Deep Learning

Artificial Intelligence: Chess & Go?



“mechanical tWolfgang von Kempelen



Deep Blue (1996)

126 mio configurations per seconds



AlphaGo (2016)

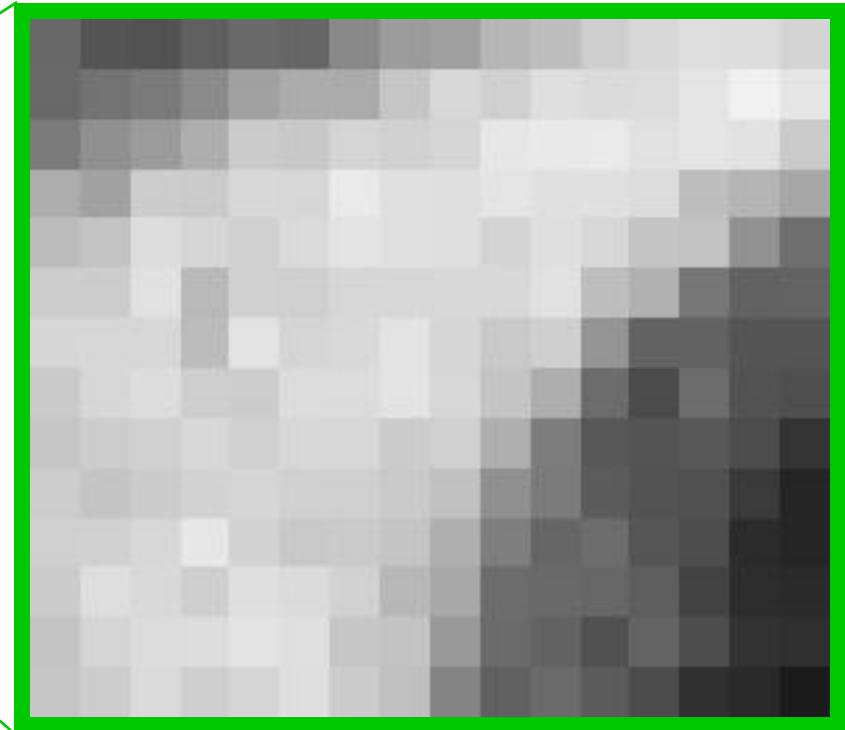
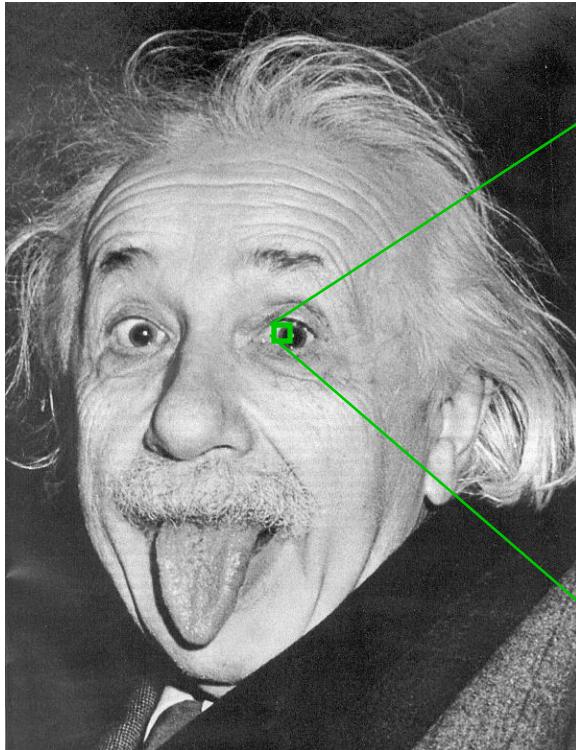
Artificial Intelligence & Computer Vision



German Museum, Munich 2024

Introduction to Deep Learning

Artificial Intelligence & Computer Vision



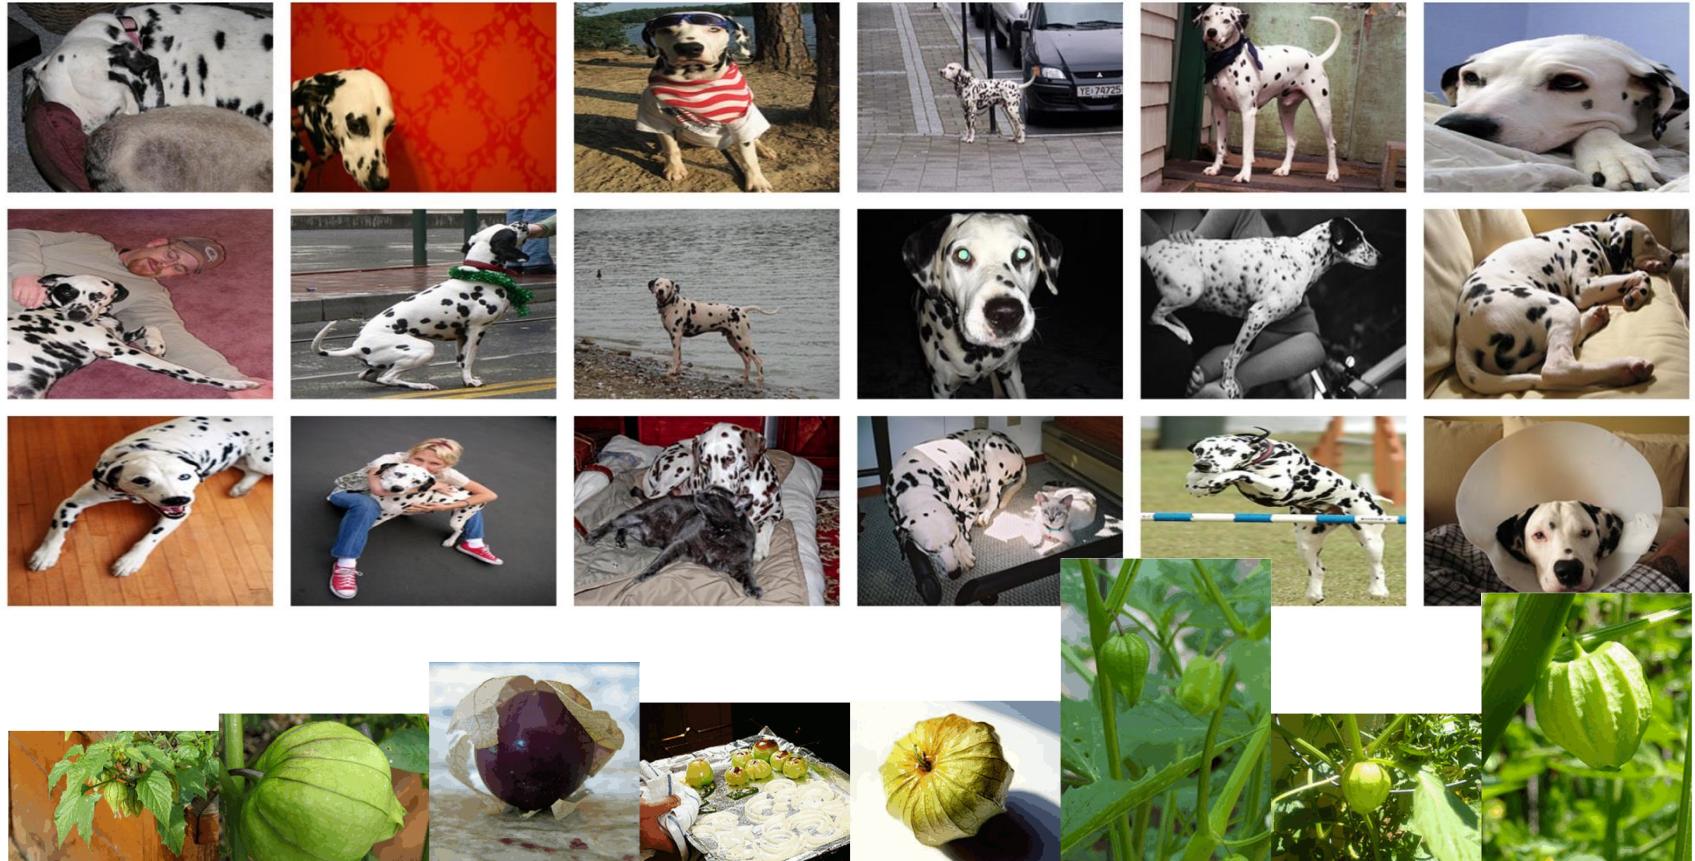
Artificial Intelligence & Computer Vision

104	85	82	96	106	102	137	156	160	182	190	206	216	222	221	212
105	115	122	139	161	172	171	198	216	208	223	219	220	229	241	229
123	144	156	174	204	200	213	209	213	231	233	234	225	229	226	203
173	162	205	203	216	215	234	224	223	229	224	224	220	190	181	167
187	196	221	214	208	218	227	224	223	212	223	215	197	194	145	110
205	205	225	185	209	208	215	215	216	216	225	190	177	119	98	100
214	215	215	188	227	213	217	227	214	201	207	149	99	98	85	85
202	215	221	206	205	220	220	227	215	197	174	108	75	108	83	79
198	205	208	216	208	216	215	204	208	175	125	88	85	88	77	52
205	197	204	211	214	209	208	202	193	144	123	92	83	81	59	38
208	209	215	231	211	201	203	197	175	127	102	109	86	78	44	38
204	222	217	207	223	219	210	184	169	108	106	103	94	67	45	43
197	213	220	221	227	224	198	195	153	107	99	81	100	78	51	47
198	205	218	207	213	223	204	192	132	97	107	91	76	48	42	27

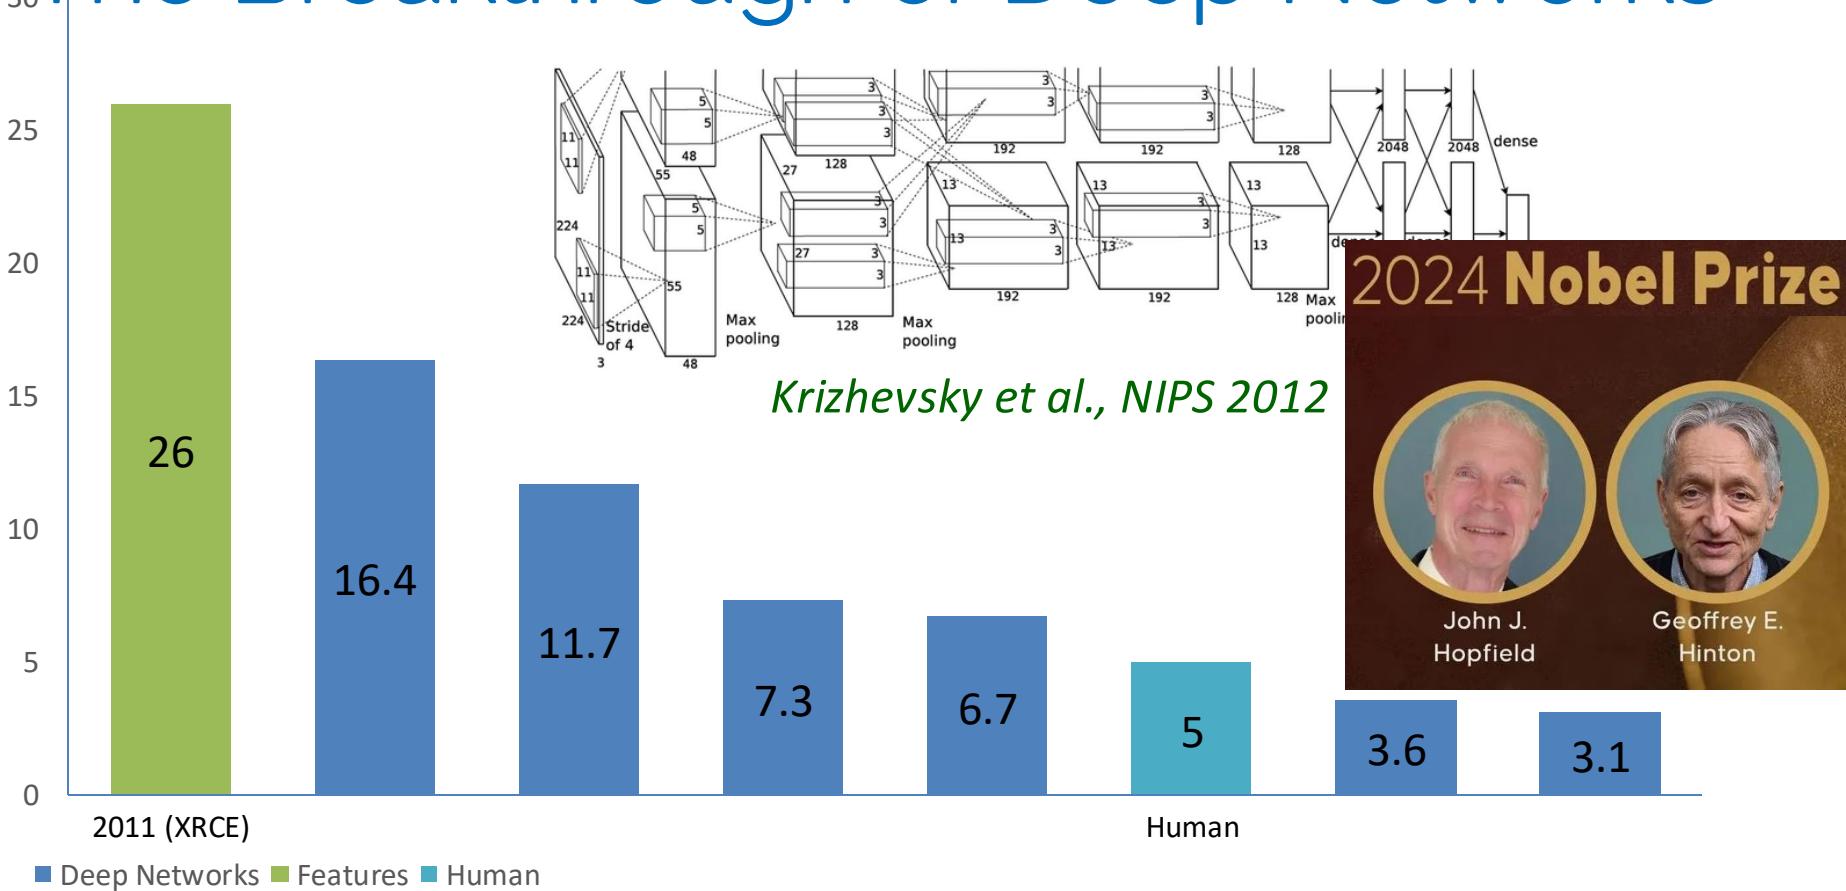


Einstein?

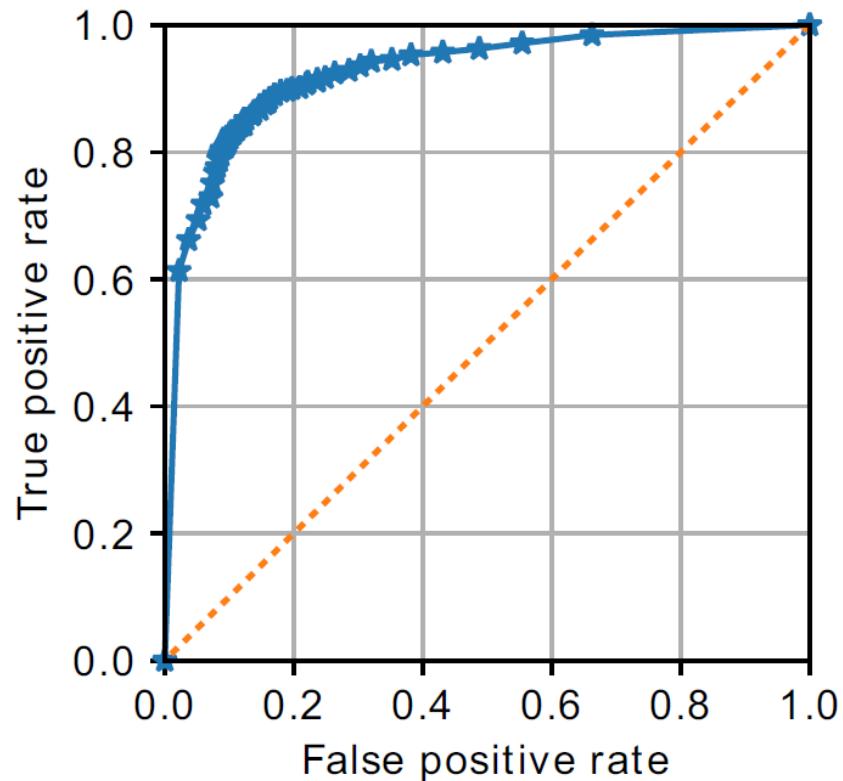
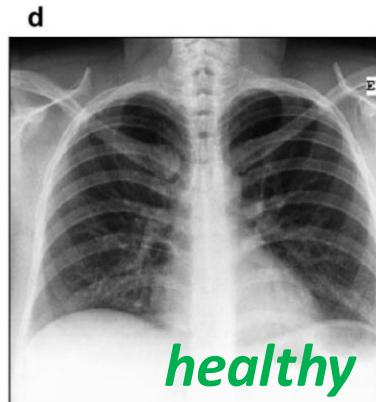
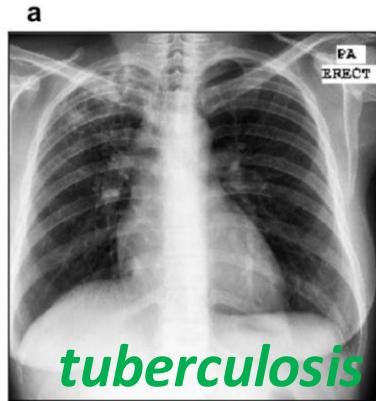
ImageNet: Objects in 14 Mio Images



The Breakthrough of Deep Networks



Deep Nets for Tuberculosis Screening



Pasa et al., "Efficient Deep Nets for X-Ray Tuberculosis Screening", Sci. Rep. 2019

Deep Nets for Protein Prediction

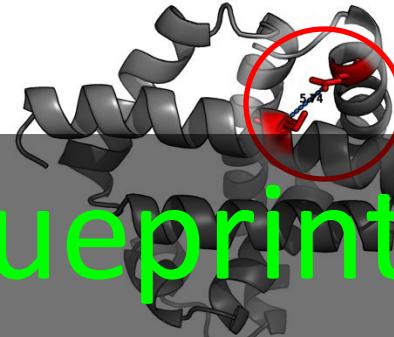
VLSGEWQLVLHVWAKVEADVAGH
GQDILIRLFKSHPETLEKFDRFKH
LKTEADMKASEDLKKHGVTVLTAL
GAILKKKGHHEAELKPLAQSHATK
HKIPIKYLEFISEAIIHVLHSRHP
GDFGADAQGAMNKALELFRKDIAA
KYKELGY (Homo sapiens)

HMM

VLSGEWQLVLHVWAKVEADVAGH
MGLSDGEWQLVNVWGKVEADIPGFKH
MGLSDGEWQLVNVWGKVEADLAGKFKTAL
HGQDVLIIRLFKGHPETLEKFDRFKHATK
HLKTEADMKASEDLKKHGNTVITAPATVLP
LGAILKKKGHHDAELKPLAQSHATSKHIAA
KHKIPIKYLEFISEAIIHVLHSRHPDMA
PAEFGADAQGAMNKALELFRKDIAA
AKYKEL (Bottlenose dolphin)

Homologous sequences

?



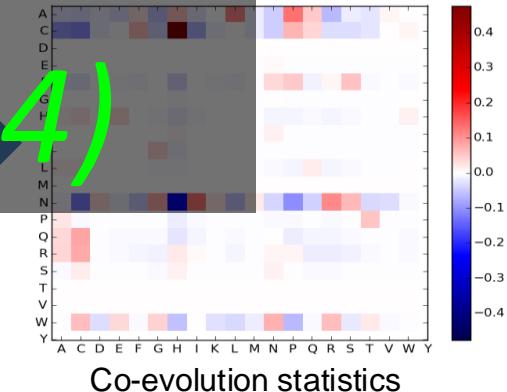
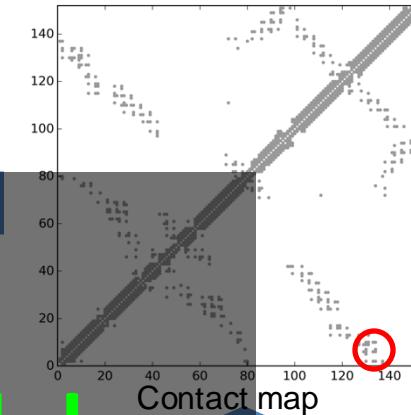
Blueprint for

Google AlphaFold
(Nobel Prize 2024)

Pairwise variants
of mutations

constraint

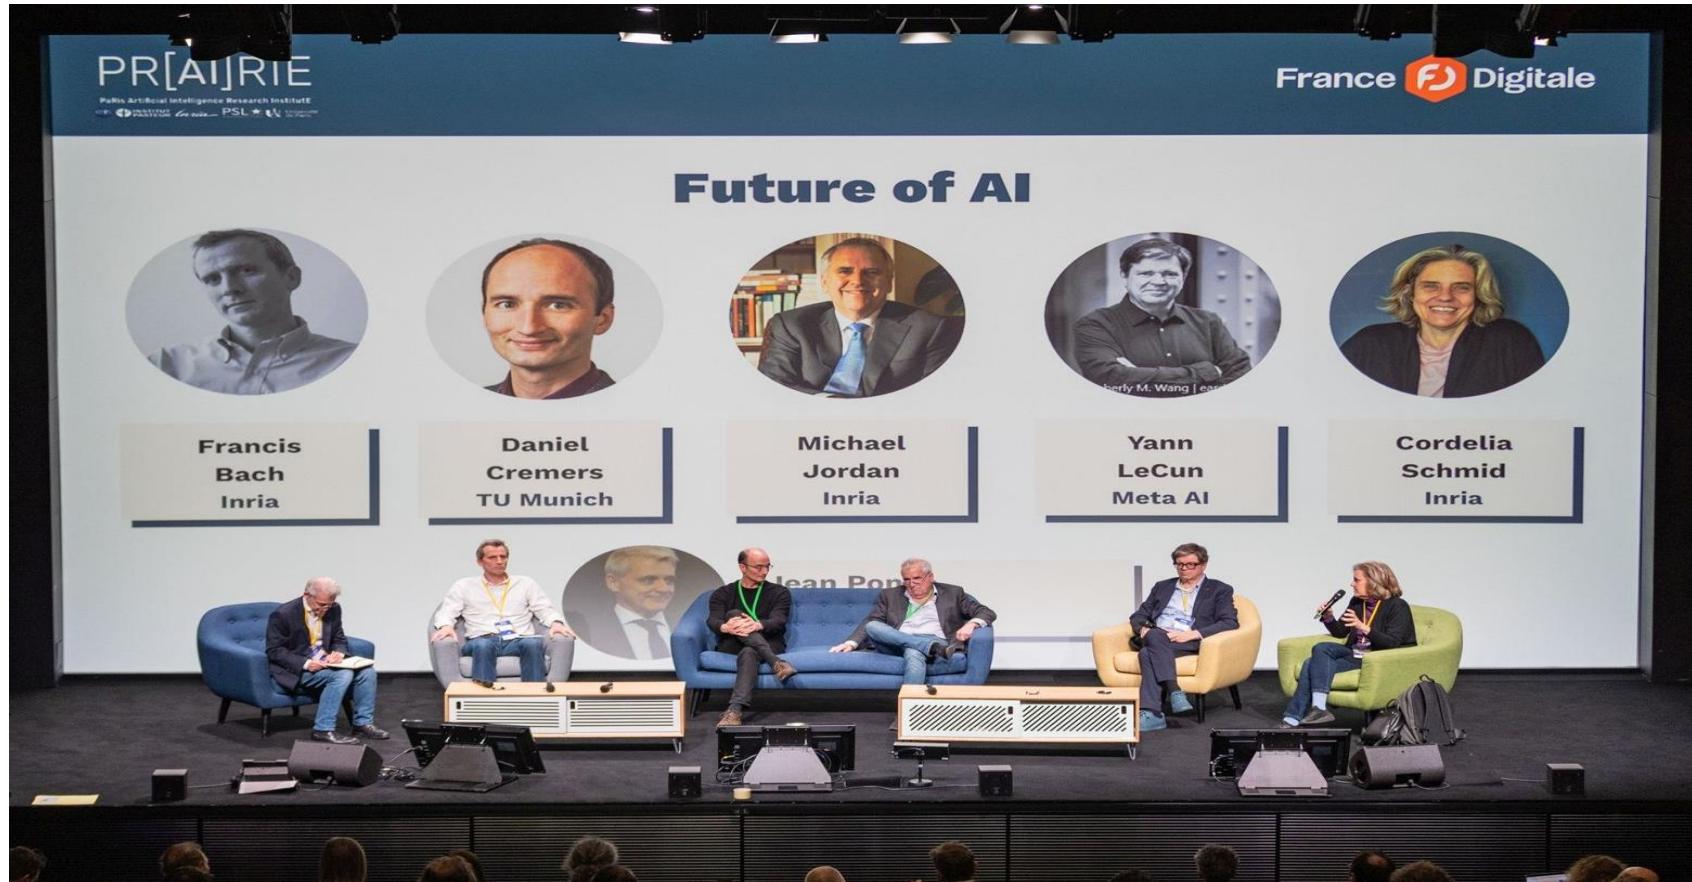
Conv
Net



Golkov,..., Cremers, "Protein Contact Map Prediction with Deep Networks", NeurIPS 2016.

Introduction to Deep Learning

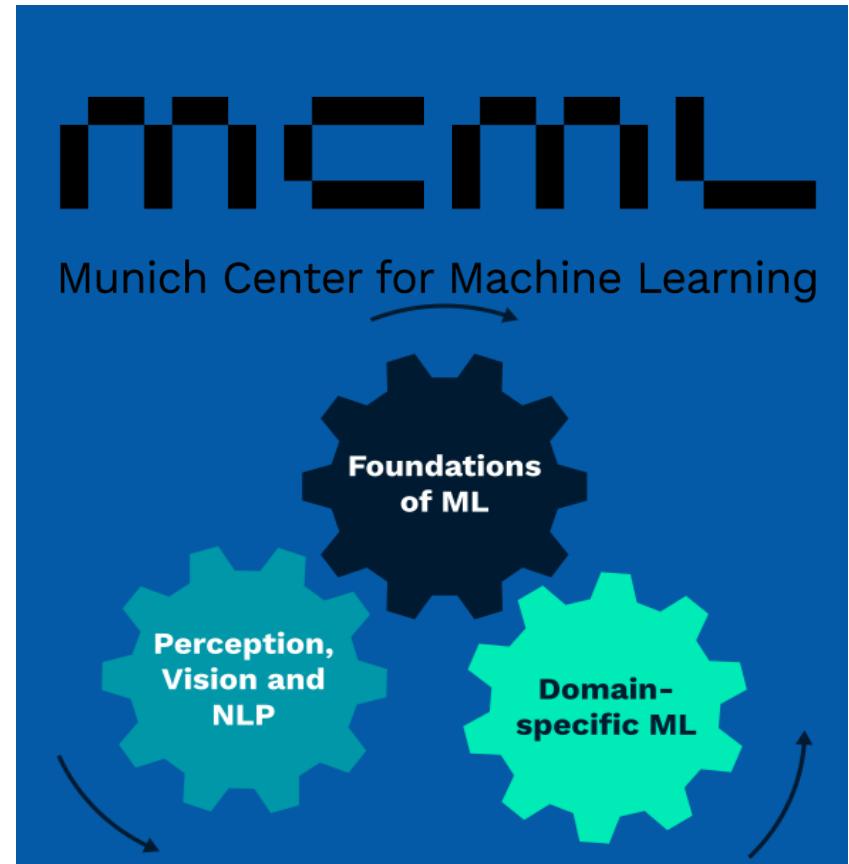
AI Panel, Paris, April 5, 2024



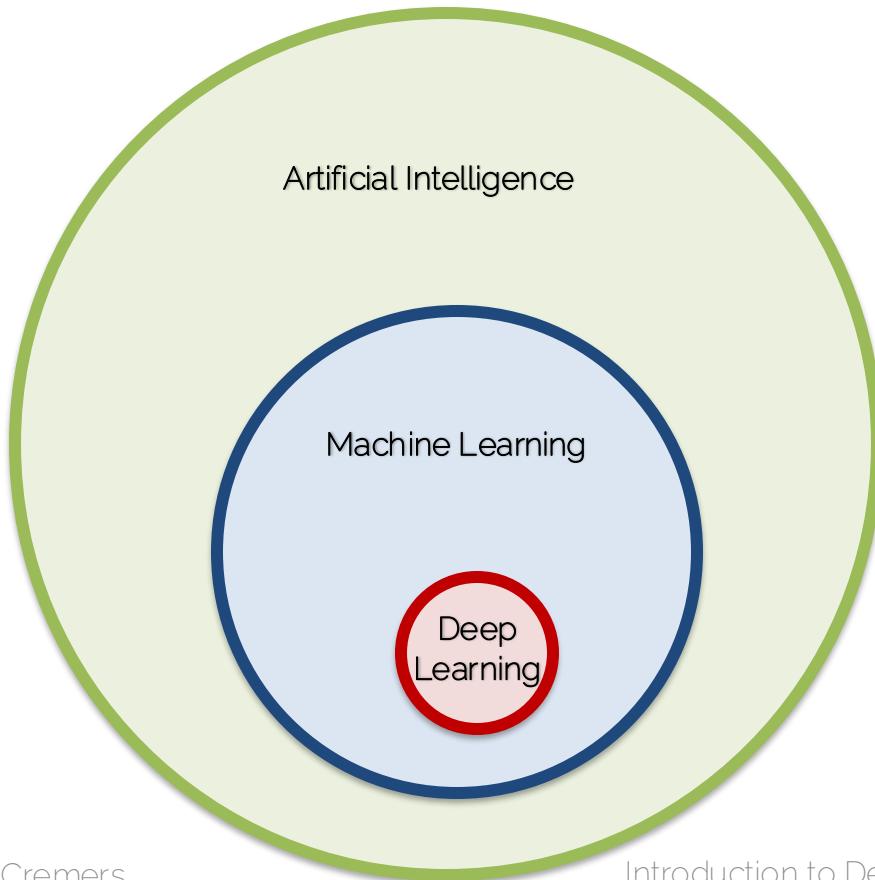
European & National Engagements



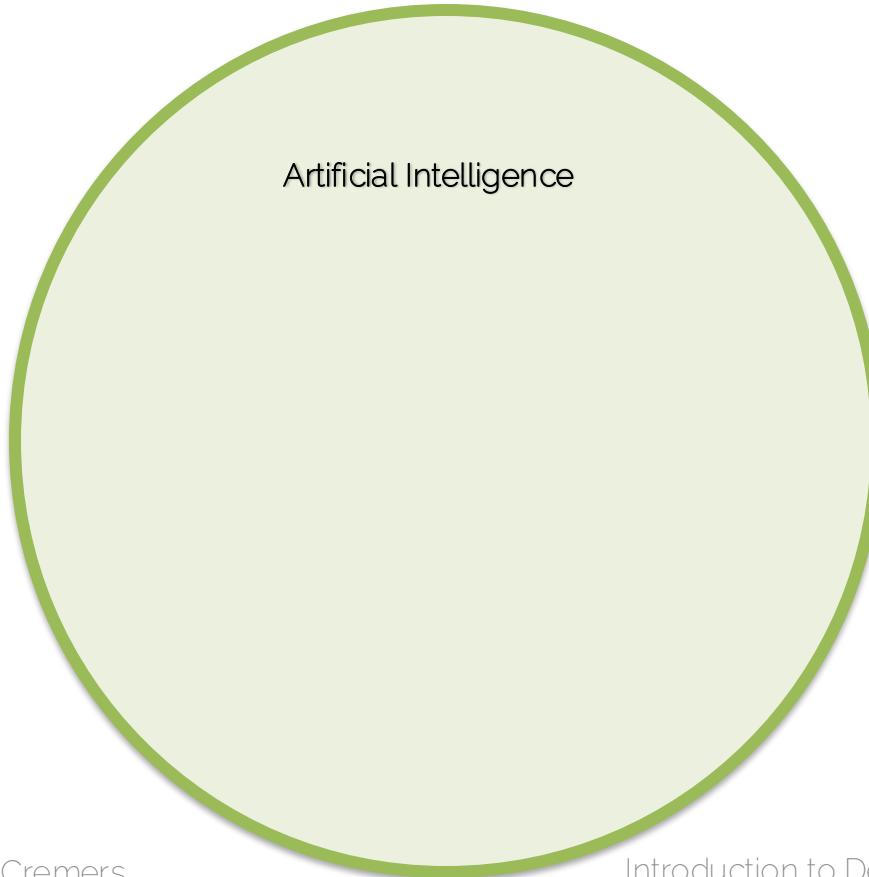
Munich Unit of the
European Lab for Learning
and Intelligent Systems (ELLIS)



What is this Lecture about?



What is this Lecture about?

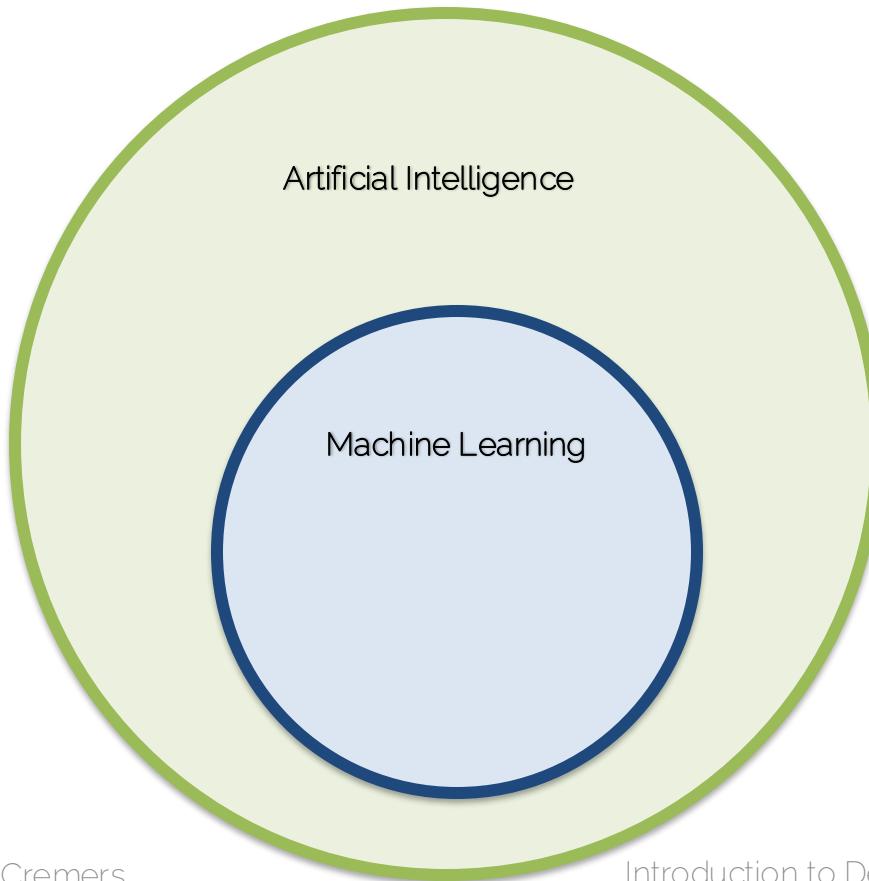


- AI Methods
 - Broad definition!
 - "if" statements

```
if (cold)
    turn_up_heat();
```
 - Binary Search
 - Dijkstra, A*, ...
 - Prime, Kurskal, ...
 - Logic algorithms, etc.

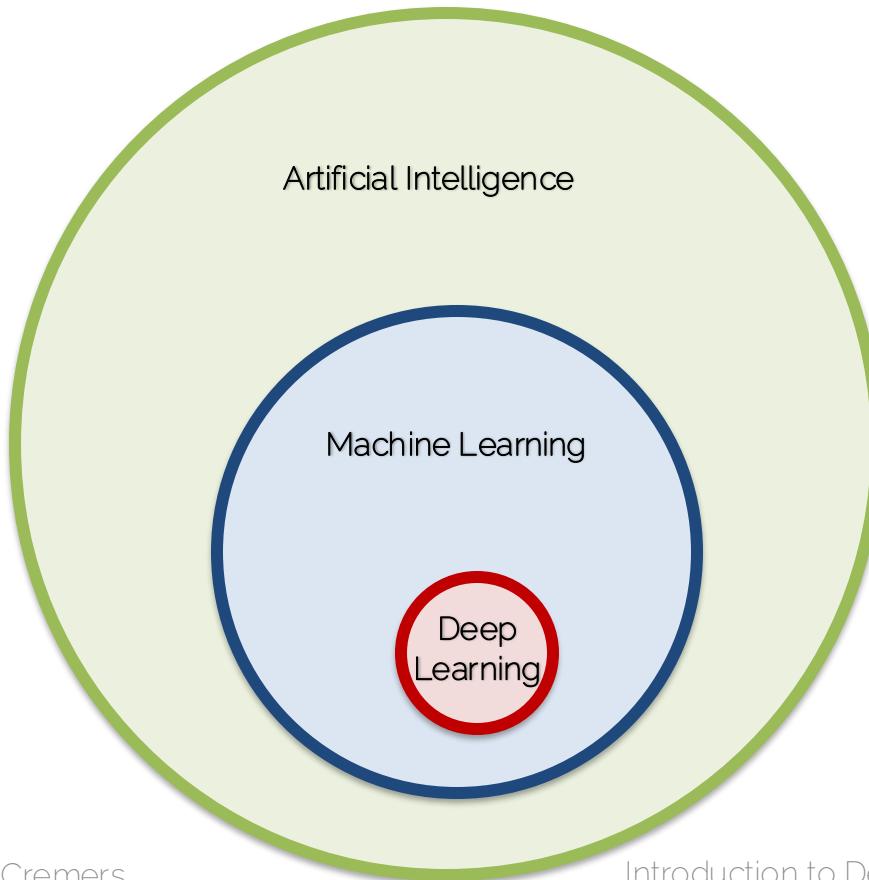
— ...

What is this Lecture about?



- ML Methods
 - Linear/logistic regression
 - Support Vector Machines
 - Random Decision Trees, Forests, Jungles, ...
 - ...

What is this Lecture about?



- Deep Learning
 - ML-methods leveraging neural networks
 - Multi-layer perceptrons
 - Convolutional neural networks
 - Recurrent neural networks
 - Transformers
 - Generative models, etc...

Application Areas

- Computer Vision
 - Medical Imaging
 - Robotics
 - Natural Language Processing (NLP)
 - Computer Graphics
- + many more ☺

What is Computer Vision?

- First defined in the 60s in artificial intelligence groups
- "Mimic the human visual system"
- Center block of robotic intelligence



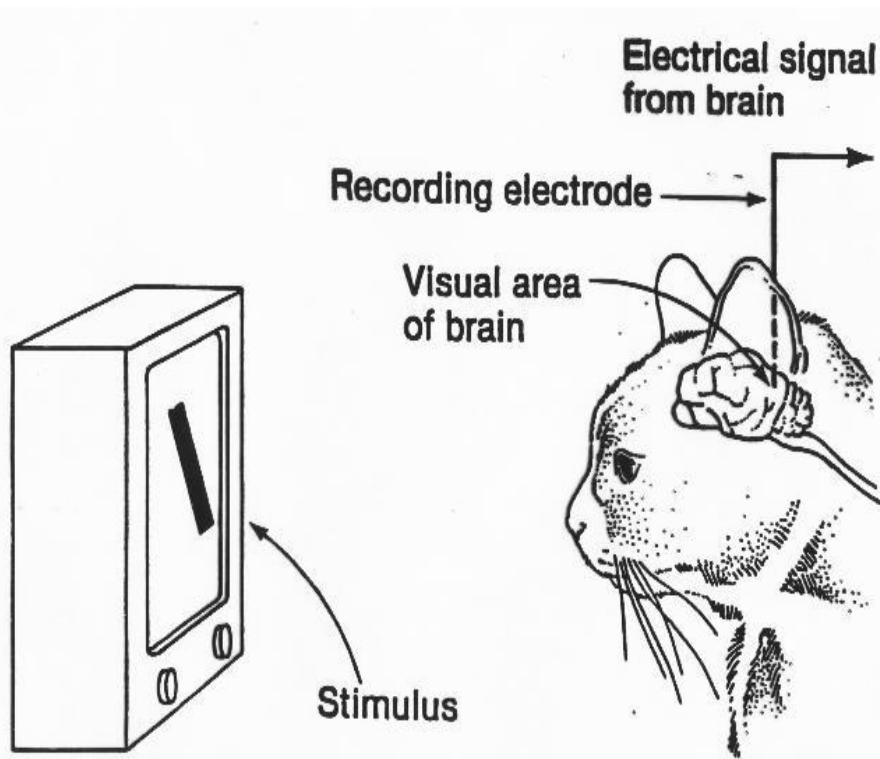
Hubel and Wiesel

- David Hubel and Torsten Wiesel were neurobiologists from Harvard Medical School
- Experiment revealed insights into the primate visual system
- Nobel prize 1981



Hubel and Wiesel Experiment

- Recorded electrical activity from individual neurons in the brains of cats.
- Slide projector to show specific patterns to the cats noted specific patterns stimulated activity in specific parts of the brain.
- Results: Visual cortex cells are sensitive to the orientation of edges but insensitive to their position



MASSACHUSETTS INSTITUTE OF TECHNOLOGY
PROJECT MAC

Artificial Intelligence Group
Vision Memo. No. 100.

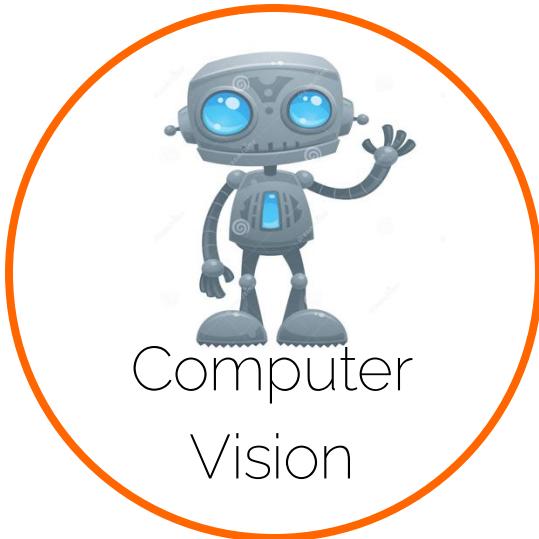
July 7, 1966

THE SUMMER VISION PROJECT

Seymour Papert

The summer vision project is an attempt to use our summer workers effectively in the construction of a significant part of a visual system. The particular task was chosen partly because it can be segmented into sub-problems which will allow individuals to work independently and yet participate in the construction of a system complex enough to be a real landmark in the development of "pattern recognition".

A Few Decades Later...



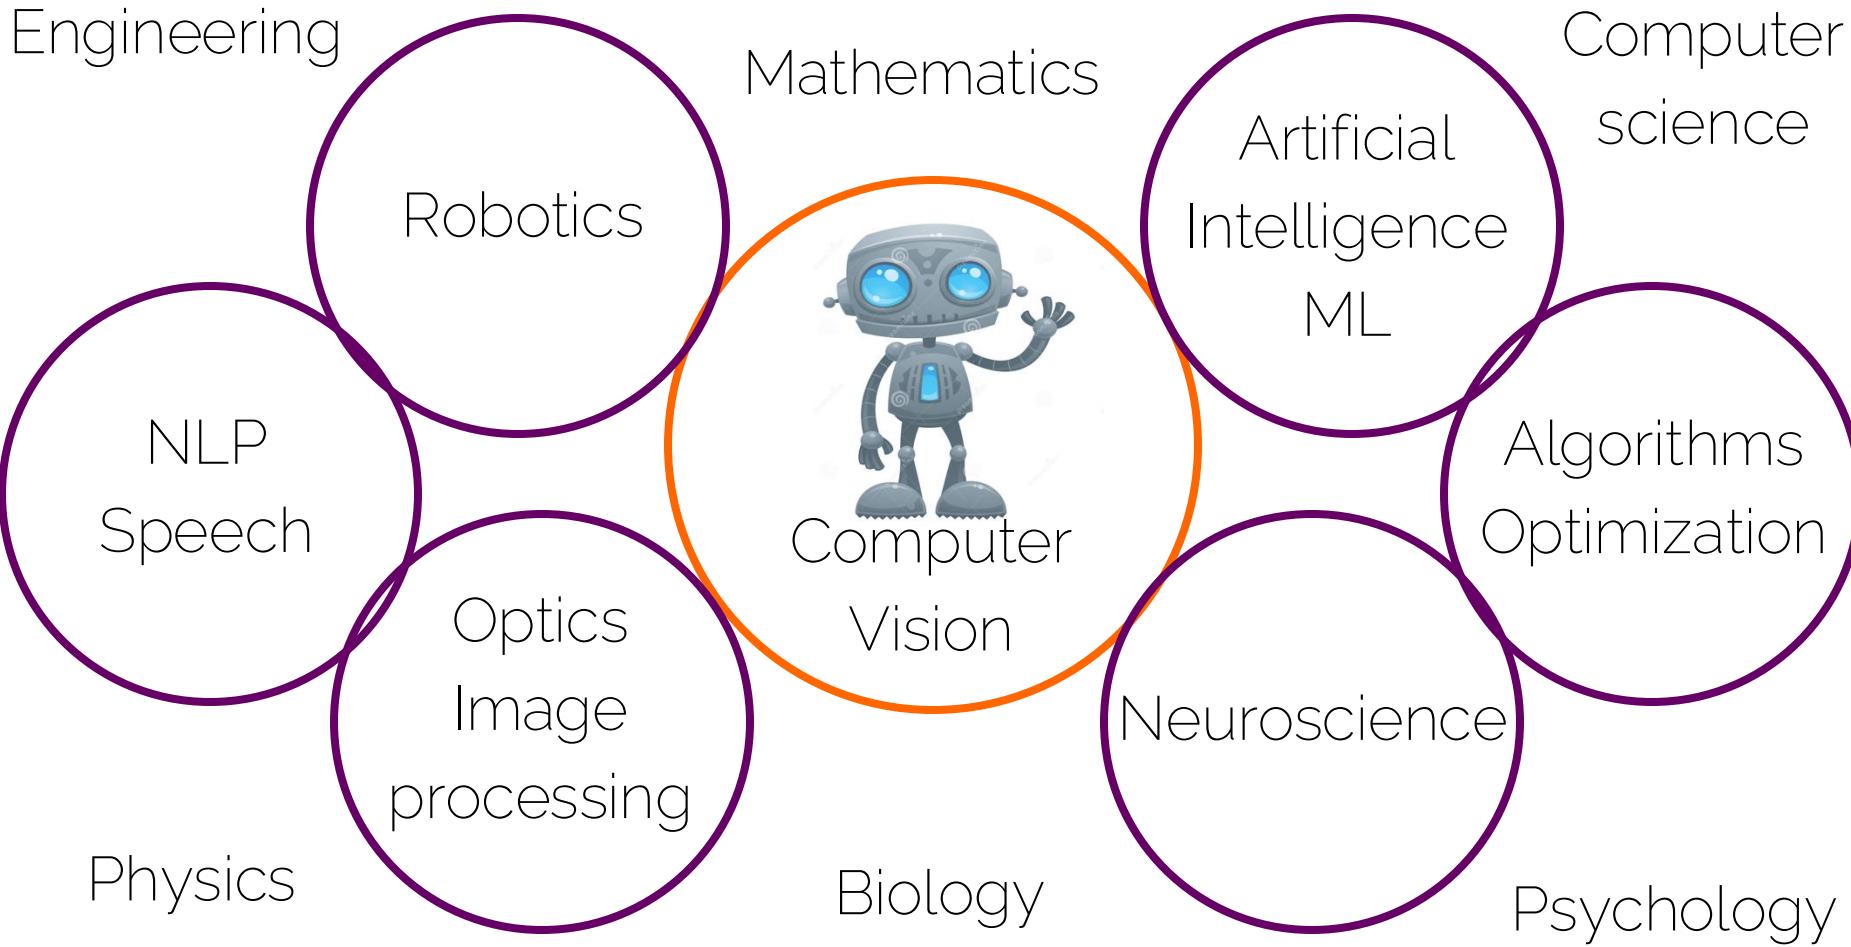


Image Classification

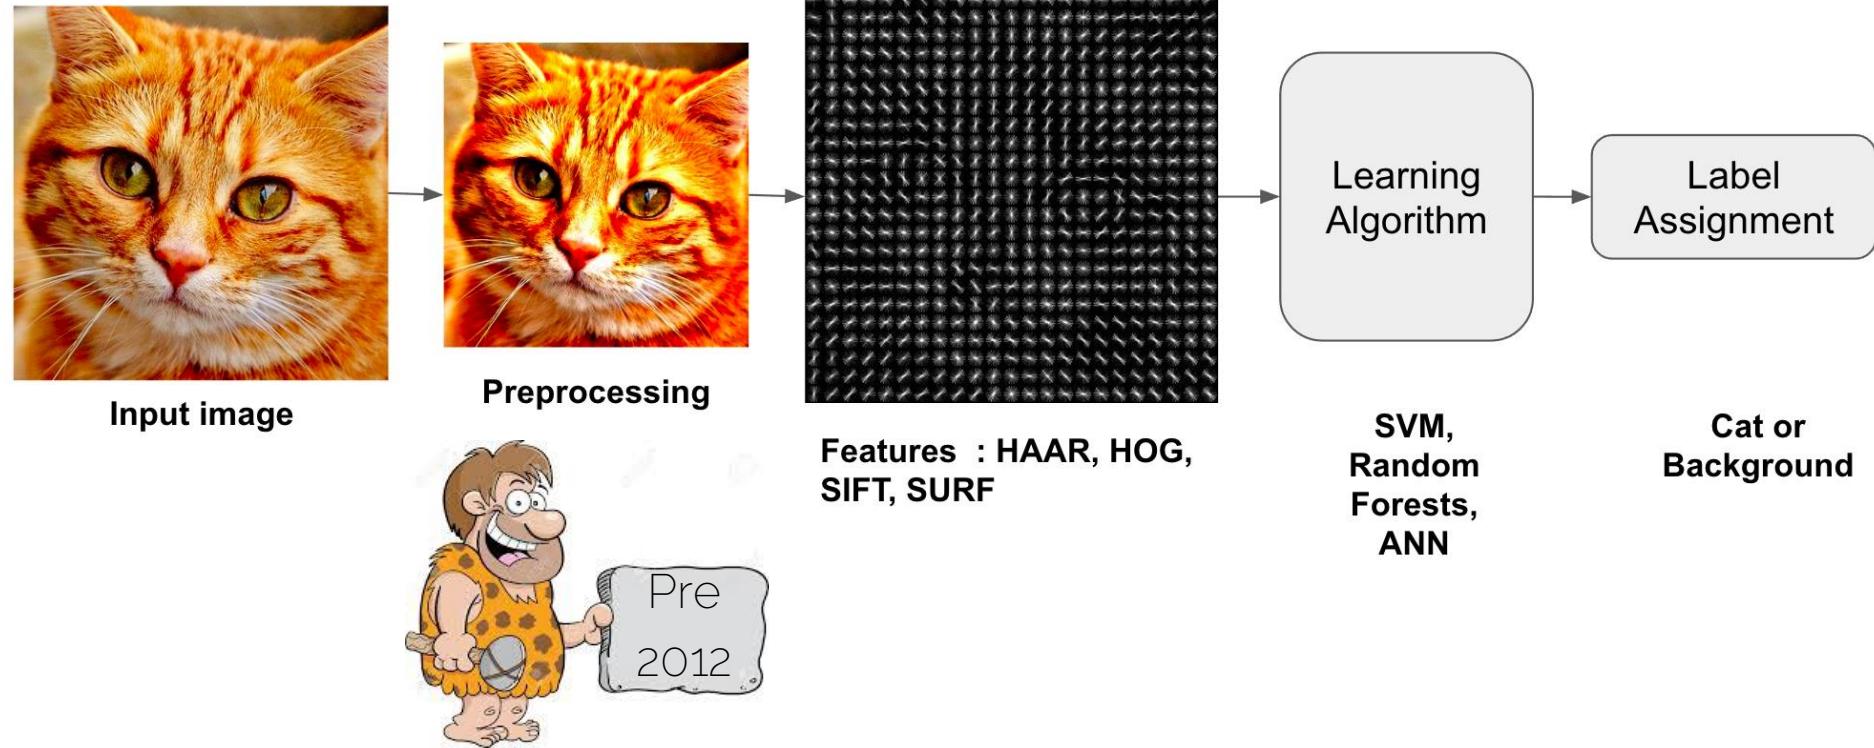


Image Classification



Input image



Awesome
magic box



Label
Assignment



Open the box

Introduction to Deep Learning

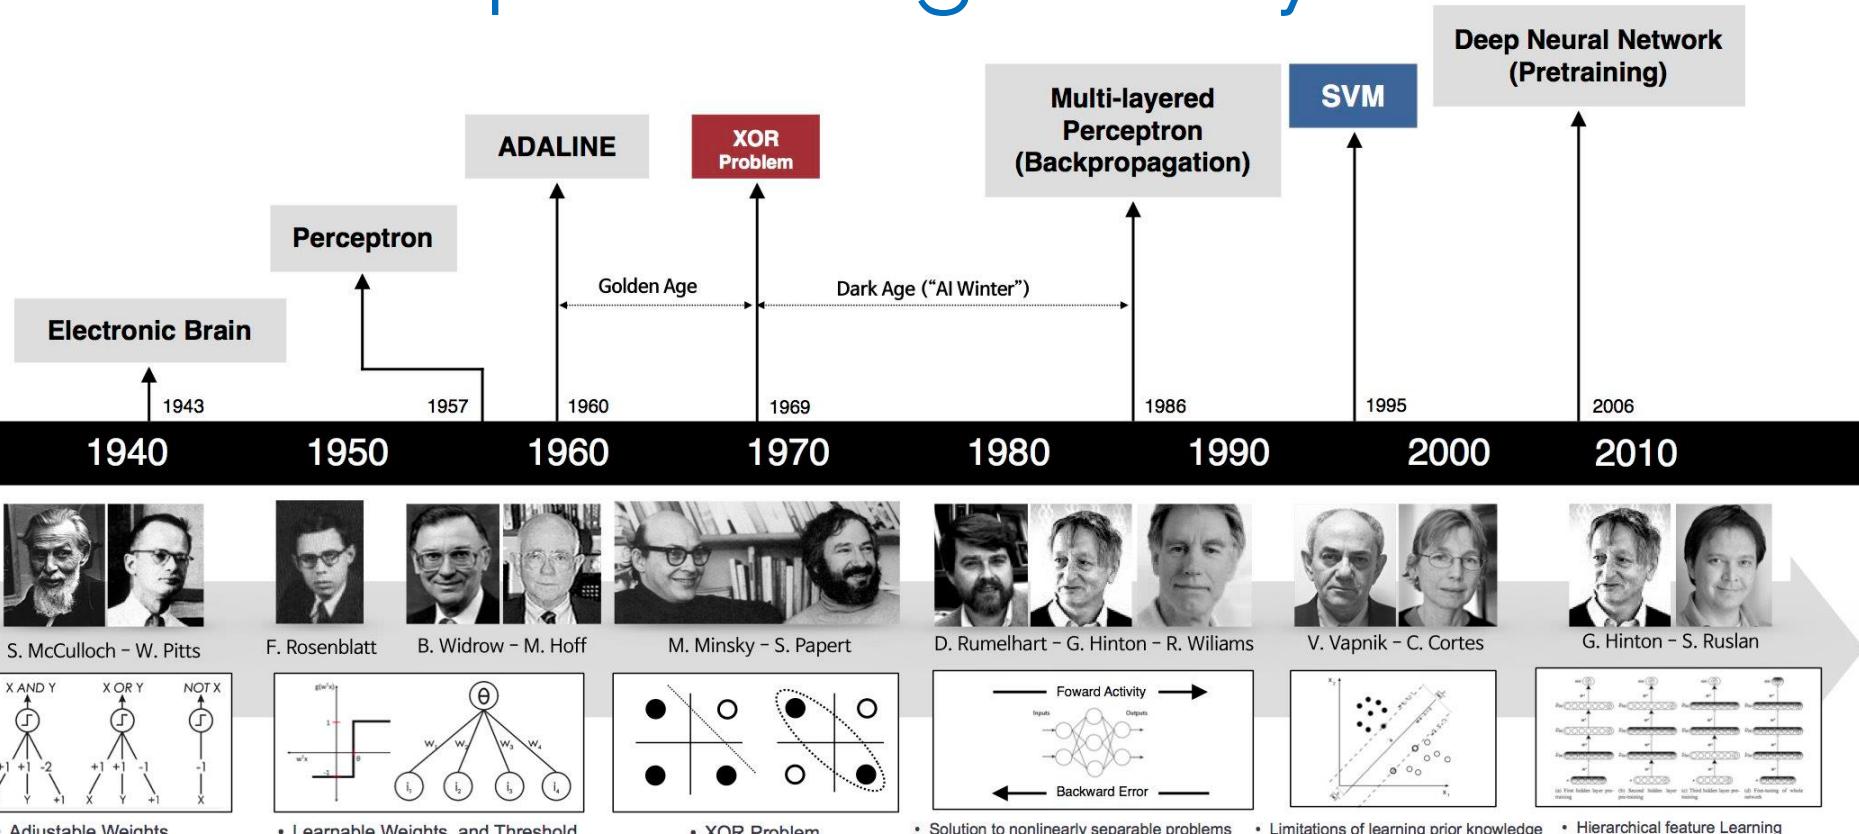


Become magicians

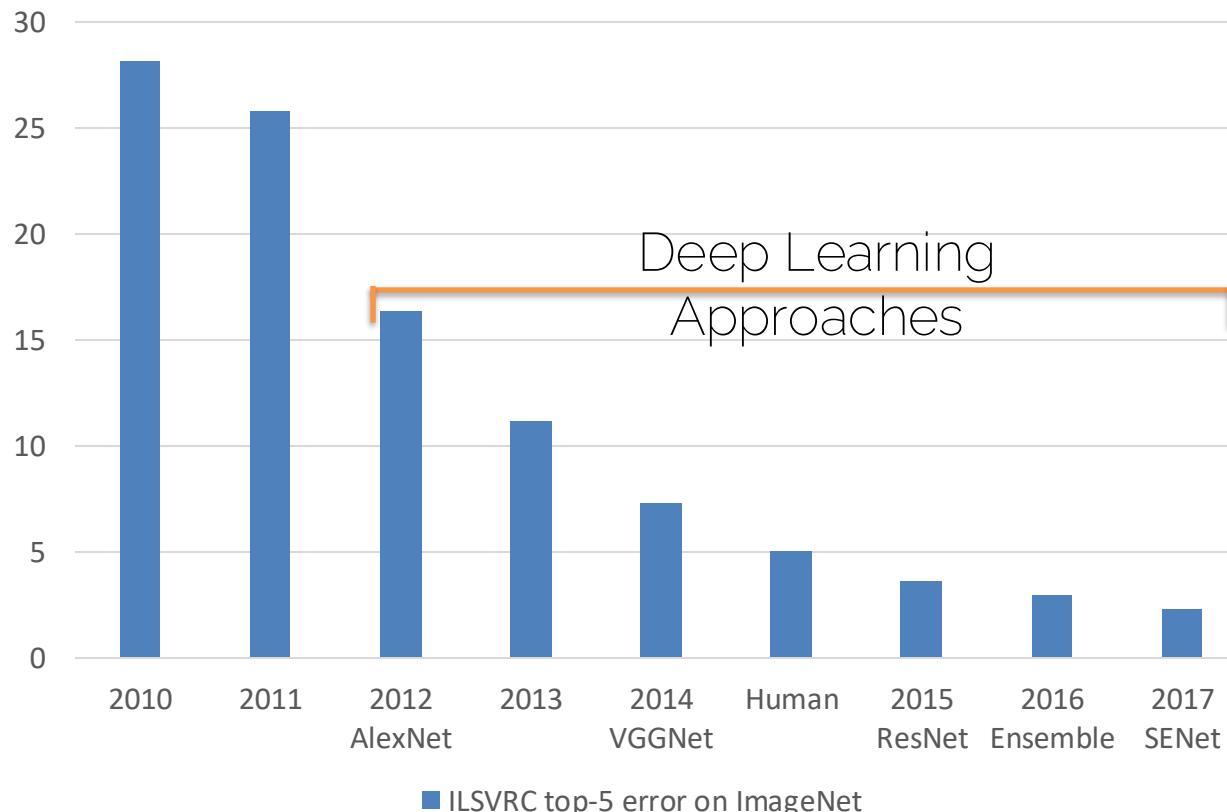
Cat or
Background

Post 2012

Deep Learning History

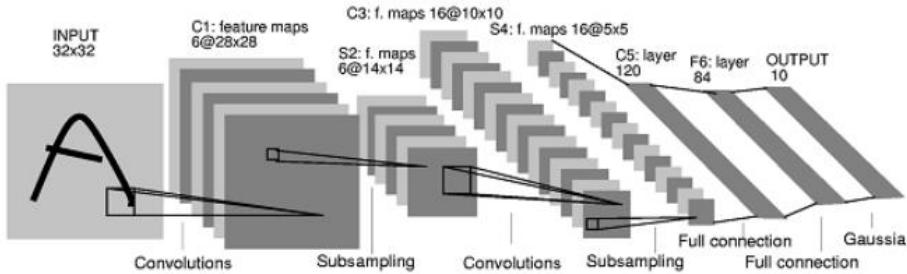


The Empire strikes Back



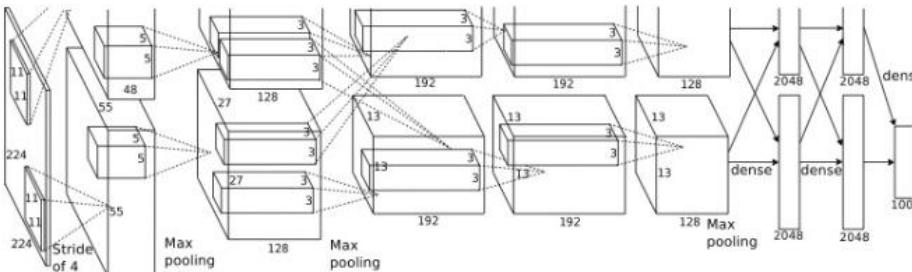
What has Changed?

1998
LeCun
et al.



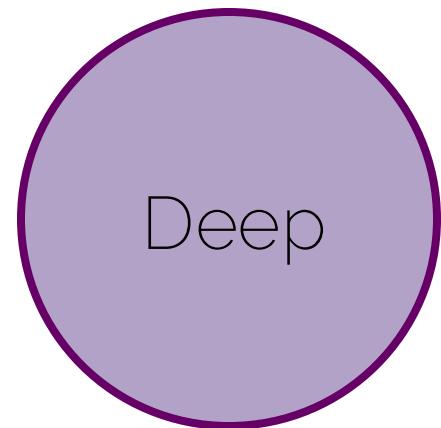
- MNIST digit recognition dataset
- 10^7 pixels used in training

2012
Krizhevsky
et al.



- ImageNet image recognition dataset
- 10^{14} pixels used in training

What Made this Possible?



Models know where
to learn from

Models are
trainable

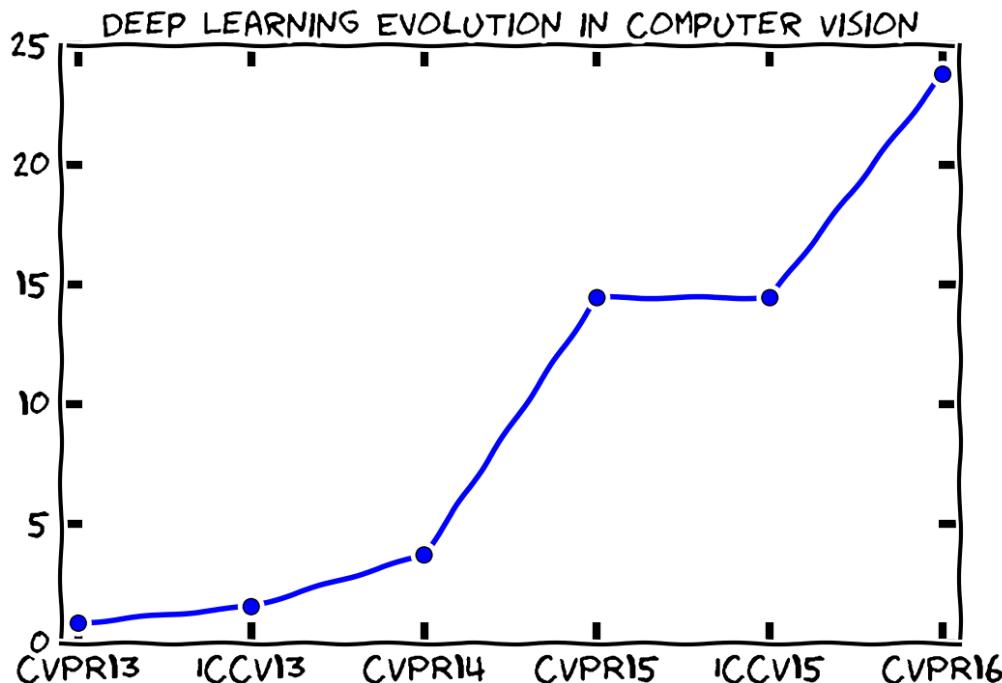
Models are
complex

Deep Learning Recognition



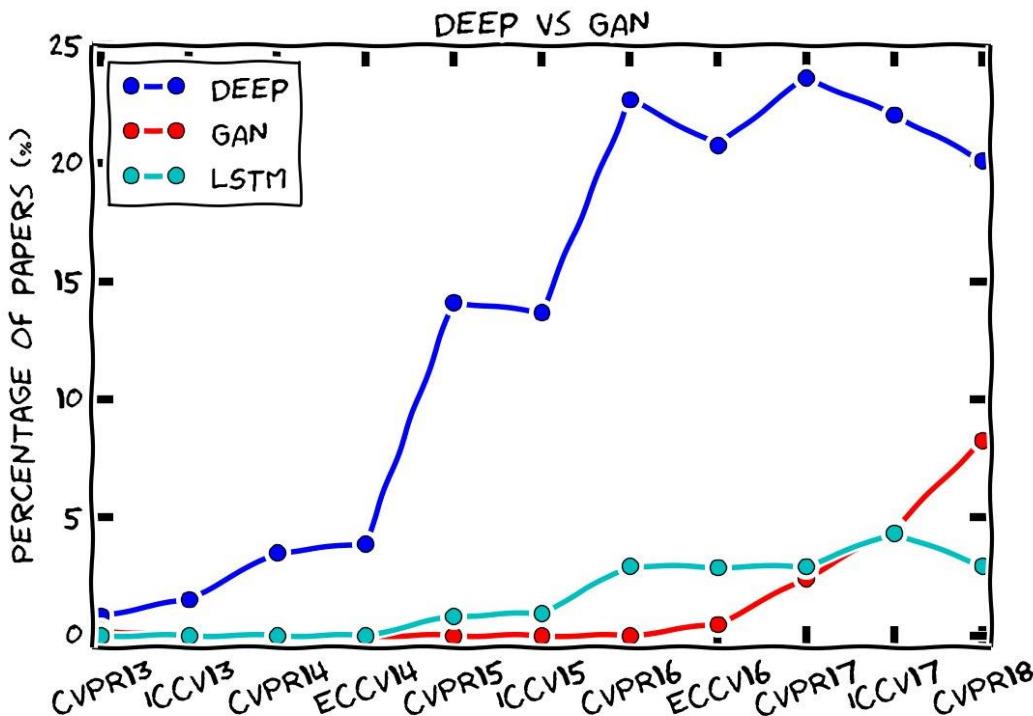
ACM Turing Award 2019 (Nobel Prize of Computing)
Yann LeCun, Geoffrey Hinton, and Yoshua Bengio

Deep Learning and Computer Vision



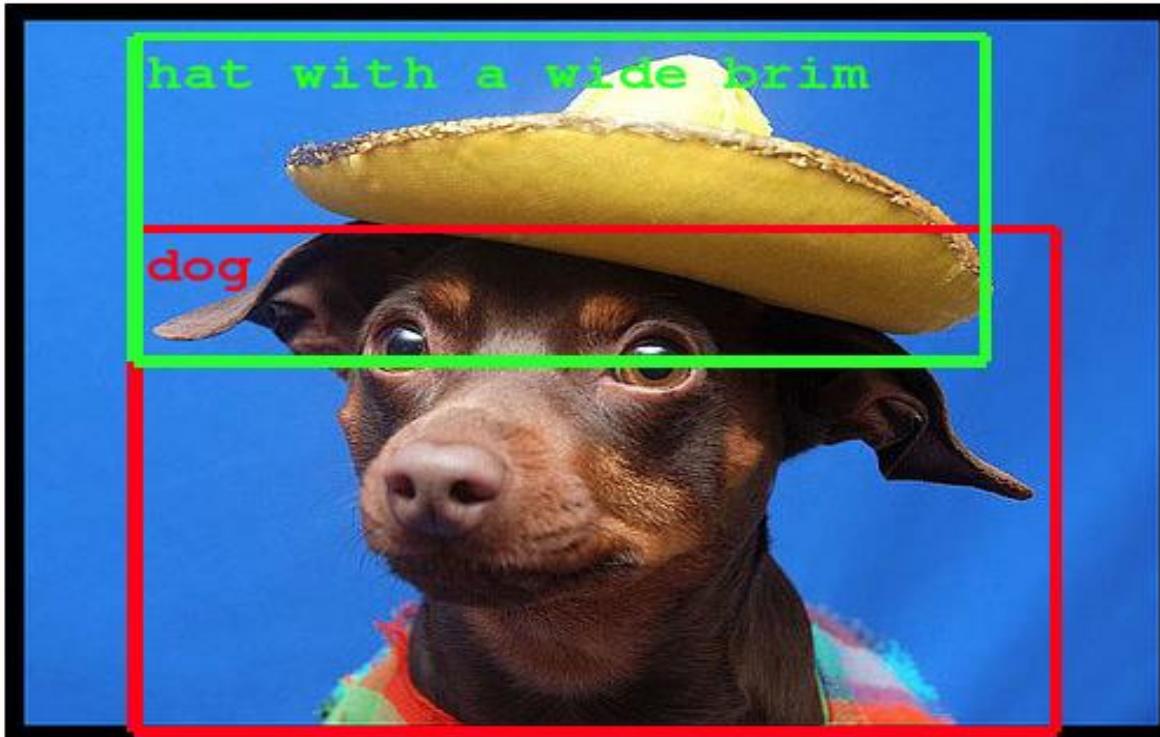
Credits: Dr. Pont-Tuset, ETH Zurich

Deep Learning and Computer Vision



Credits: Dr. Pont-Tuset, ETH Zurich

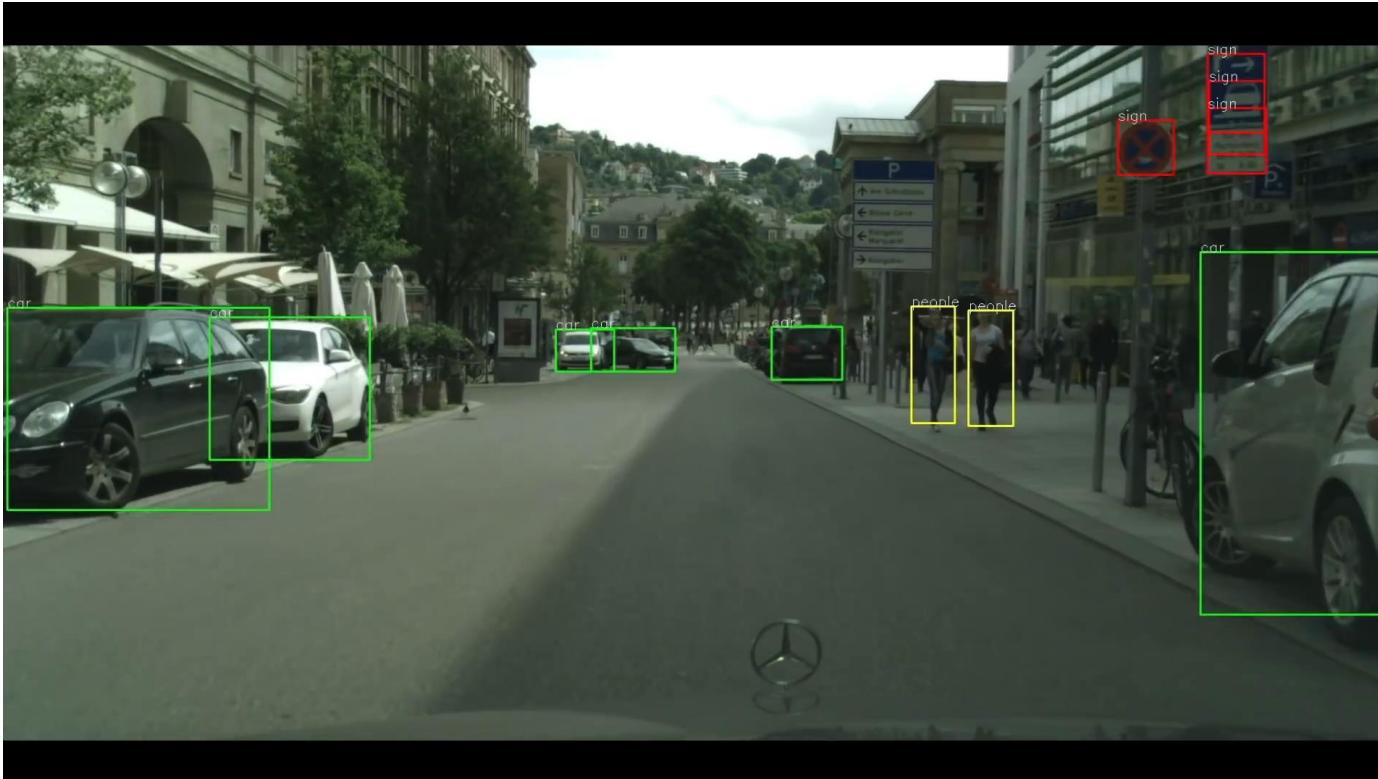
Deep Learning Today



Object Detection

Introduction to Deep Learning

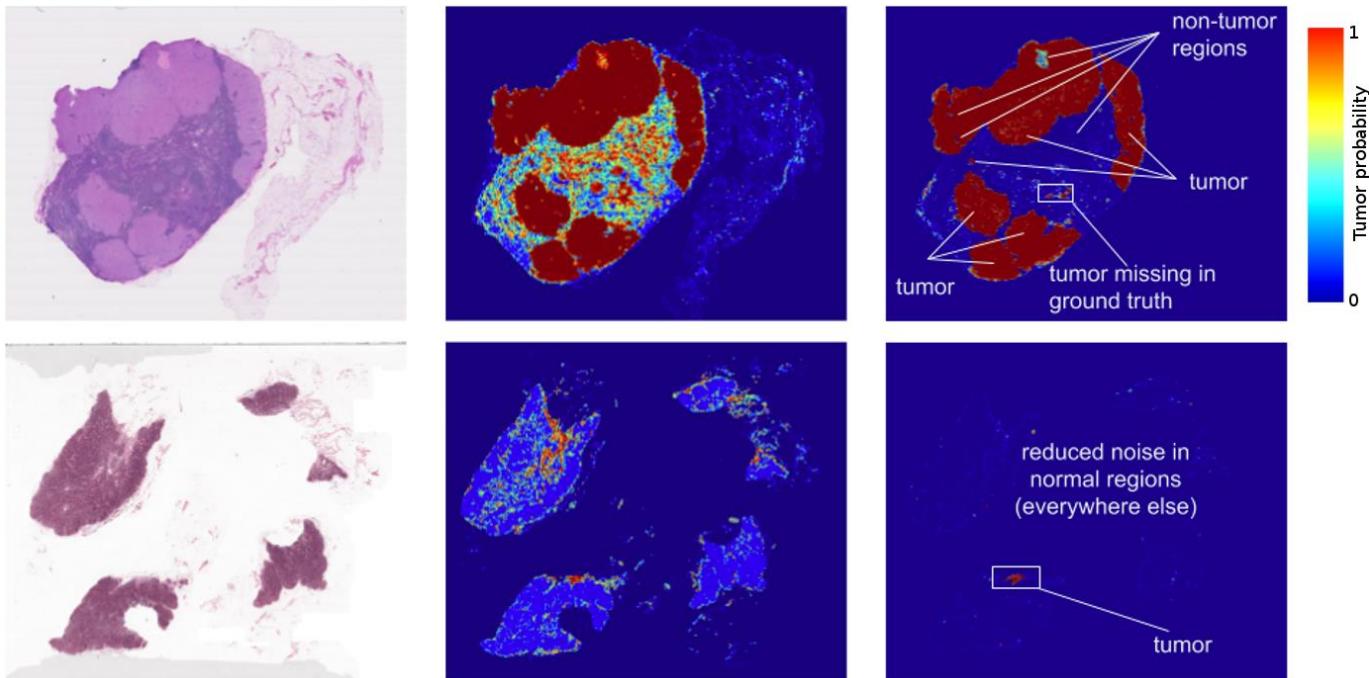
Deep Learning Today



Self-driving cars

Introduction to Deep Learning

Deep Learning Today



Healthcare, cancer detection

Deep Learning Today



AlphaGo

ever punch a cactus?



Emoticon suggestion

English – detected ▾

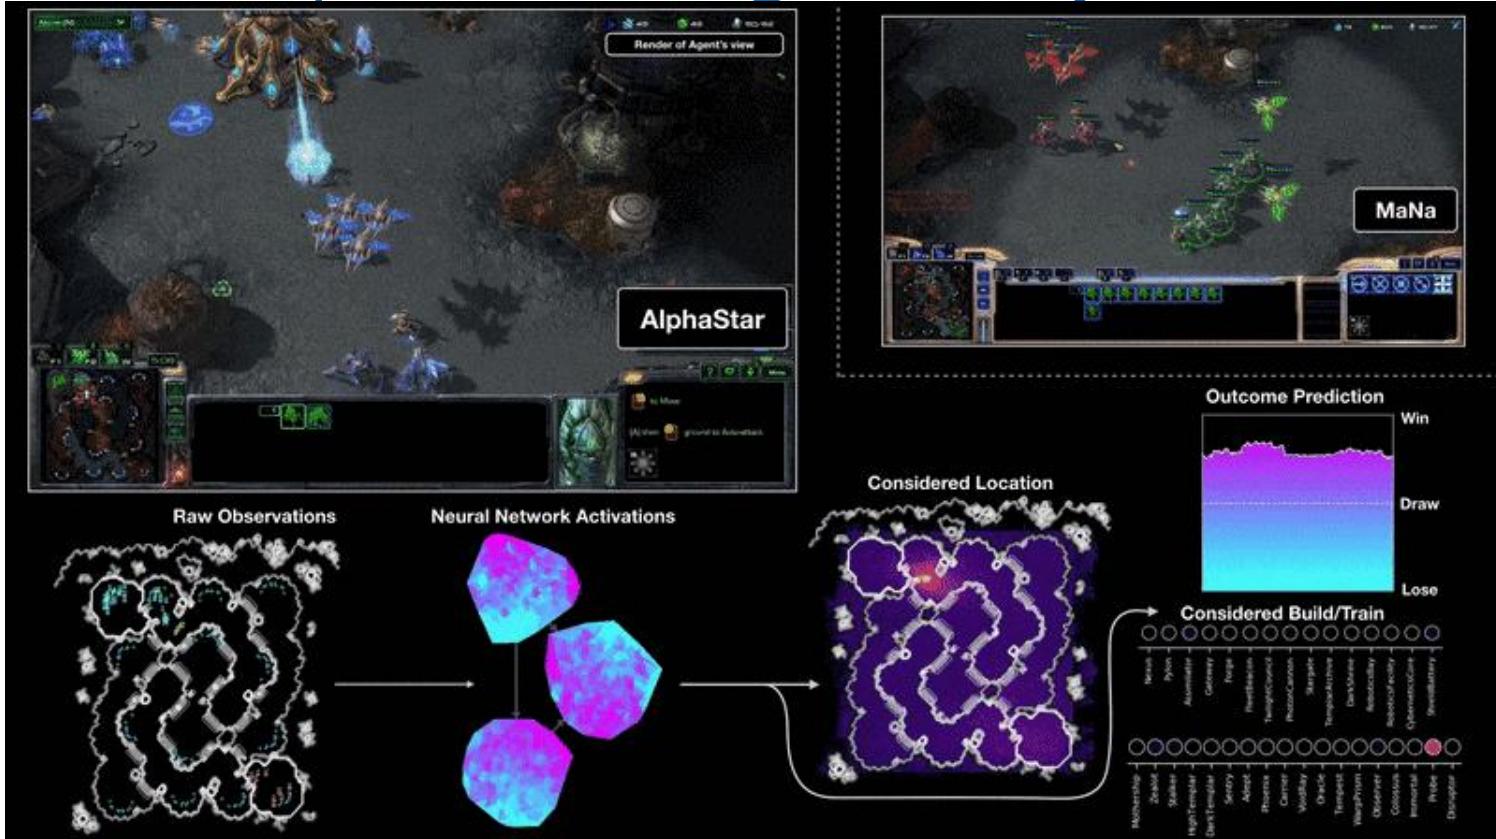
German ▾

Deep Learning rocks

Deep Learning Felsen

Machine translation

Deep Learning Today



Alpha Star
Introduction to Deep Learning

Deep Learning Today



Google LaMDA (Google IO'22)

Deep Learning Today

- Large Language Models (LLMs), e.g., Chat-GPT

NI

Tell me about deep learning and how I can teach it at TUM.



Deep learning is a subset of machine learning that involves training artificial neural networks to recognize patterns and make predictions or decisions based on input data. The term "deep" refers to the fact that these networks have many layers, which allows them to learn increasingly complex representations of the data as it moves through the network. thumb up thumb down

To teach deep learning at TUM, you could start with an introductory course that covers the basics of machine learning and neural networks. This could include topics such as:

1. Supervised learning: This involves training a model using labeled examples, so that it can make predictions on new, unseen data.
2. Unsupervised learning: This involves training a model to identify patterns in data without explicit labels.

Deep Learning Today

TEXT DESCRIPTION

An astronaut **Teddy bears** A bowl of soup

mixing sparkling chemicals as mad scientists **shopping for groceries** working on new AI research

in the style of ukiyo-e as a one-line drawing in ancient Egypt



DALL-E 2



Dall-E 2 Image generation from text

Deep Learning Today

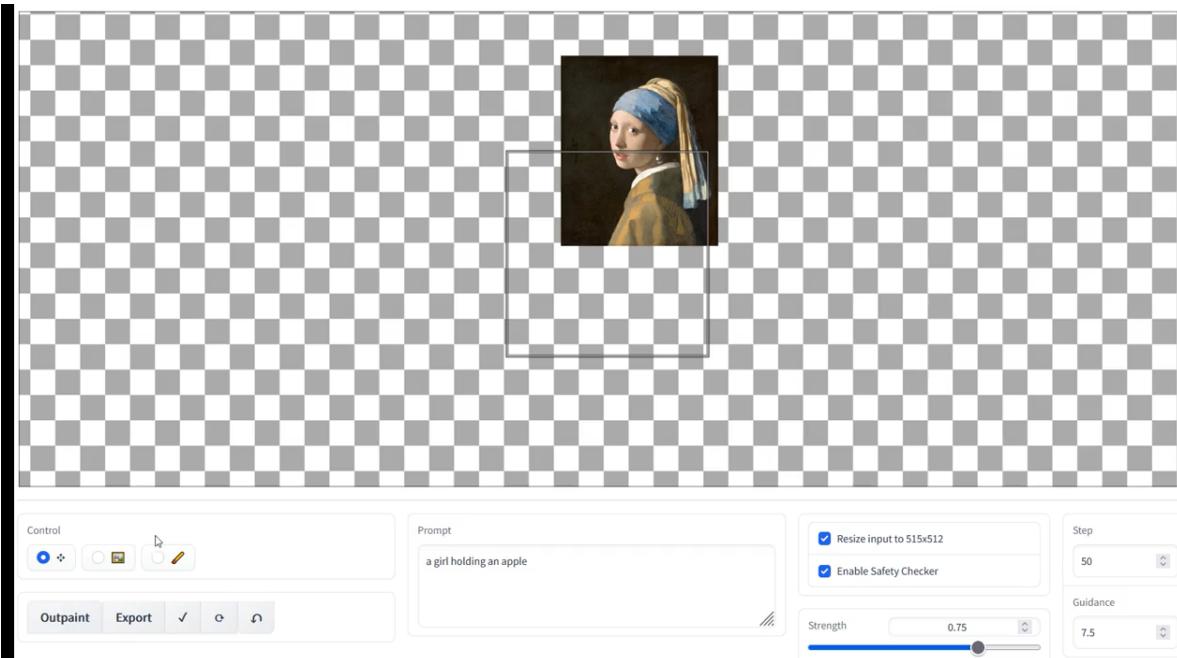
Darth Vader
on a unicorn
in Oxford



Esser, Rombach, Ommer CVPR 2021 Rombach et al, CVPR 2022

Introduction to Deep Learning

Deep Learning Today



StableDiffusion Image Outpainting

Deep Learning Today



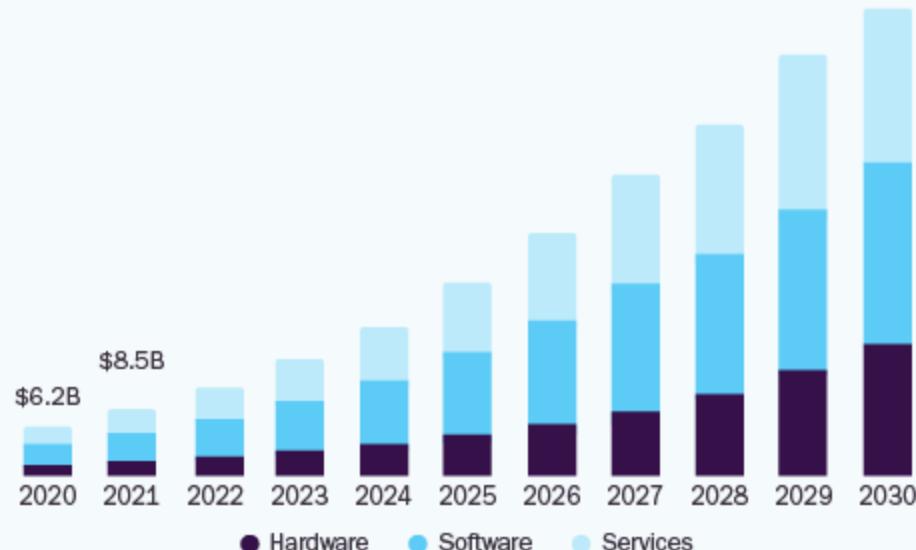
Deep Learning Today



Deep Learning Market

U.S. Deep Learning Market

size, by solution, 2020 - 2030 (USD Billion)



GRAND VIEW RESEARCH

24.1%

U.S. Market CAGR,
2022 - 2030

Source:
www.grandviewresearch.com

[...] market re
market is expected to be worth USD 415 Billion by 2030.

Deep Learning Job Perspective

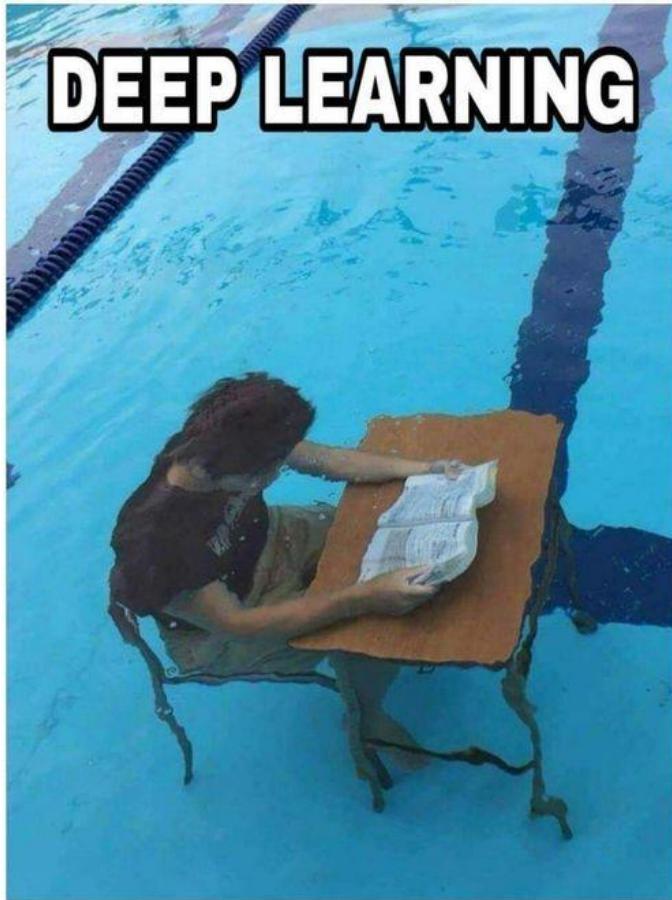
- Excellent Job Perspectives!
 - Automation requires ML/DL -> growth!
 - Top-notch companies will gladly hire you!
- Many industries now:
 - IT-Companies
 - Cars, Logistics, Health Care, etc...
 - Manufacturing / Robotics, etc...

But: Also Challenging!

- High-level understanding is not enough
 - Need proper theory background
 - Need proper practical skillsets
- Can be competitive!
 - Many good people
 - Downloading scripts / running code not enough ☺
 - Deeper understanding often requires PhDs

Deep Learning on the Internet

Deep Learning Memes

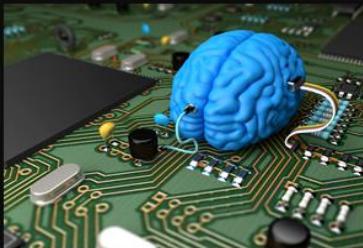


Deep Learning Memes

Deep Learning



What society thinks I do



What my friends think I do



What other computer
scientists think I do



What mathematicians think I do



What I think I do

```
from theano import *
```

What I actually do

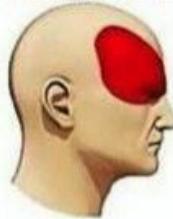
Deep Learning Memes



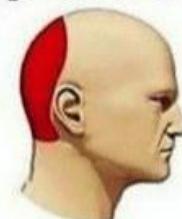
Deep Learning Memes

Types of Headaches

Migraine



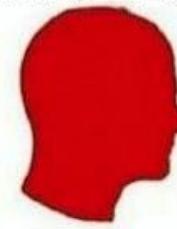
Hypertension



Stress



MATH BEHIND DL



imgflip.com

Deep Learning at TUM

Many TUM Research Labs use DL

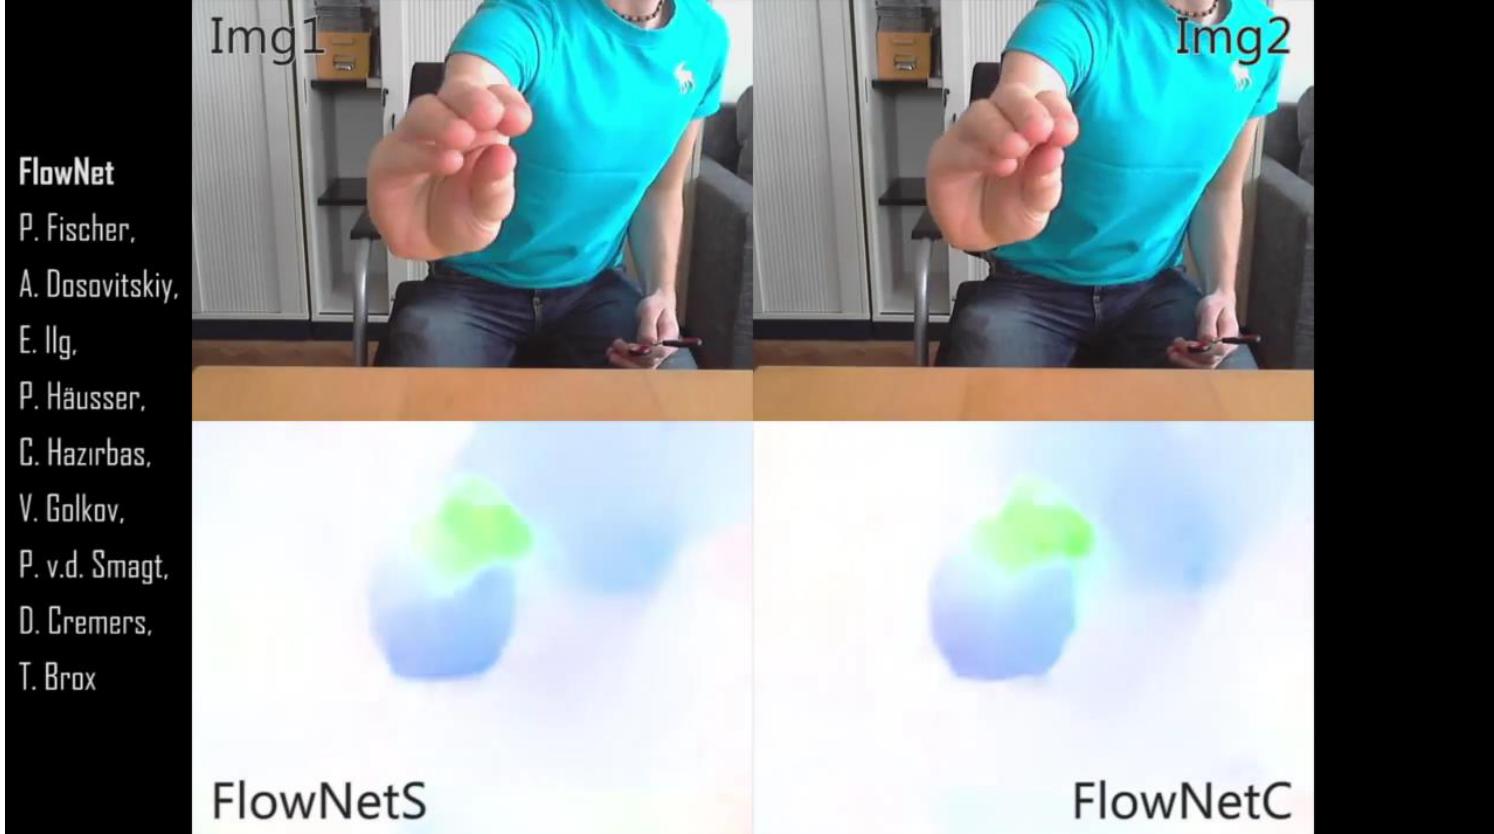
- Computer Vision Group (Prof. Cremers)
 - Research in computer vision, machine learning and robotics
- Visual Computing Lab (Prof. Niessner):
 - Research in computer vision, graphics, and machine learning
- 3D AI Lab (Prof. Dai)
 - Research in 3D perception, 3D scene understanding
- Data Mining and Analytics Lab (Prof. Günnemann)
 - Research methods for robust machine learning
- Computer Aided Medical Procedures (Prof. Navab)
 - Research in machine learning for medical applications
- And many more ☺

Deep Learning at TUM



[Caelles et al., CVPR' 17] One-Shot Video Object Segmentation

Deep Learning at TUM

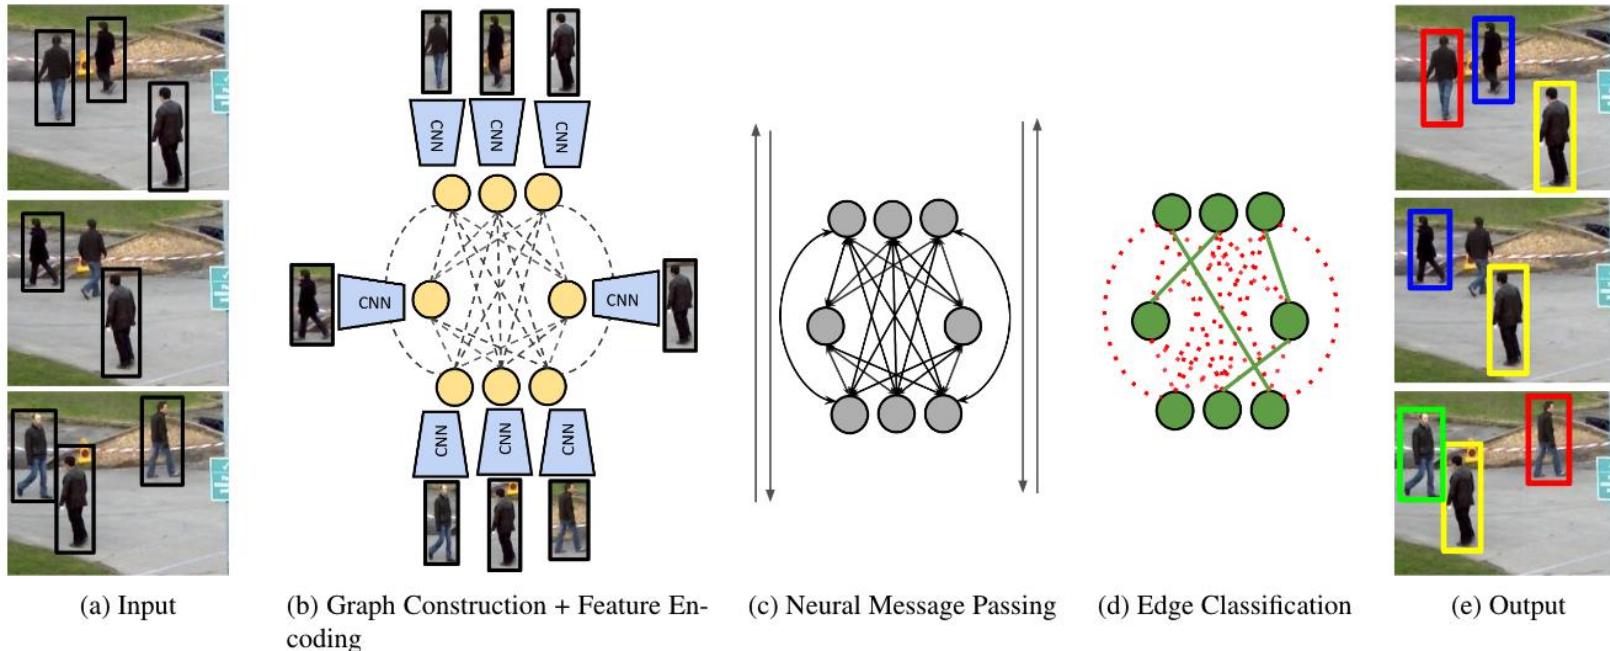


[Dosovitskiy et al., ICCV' 15] FlowNet

Introduction to Deep Learning

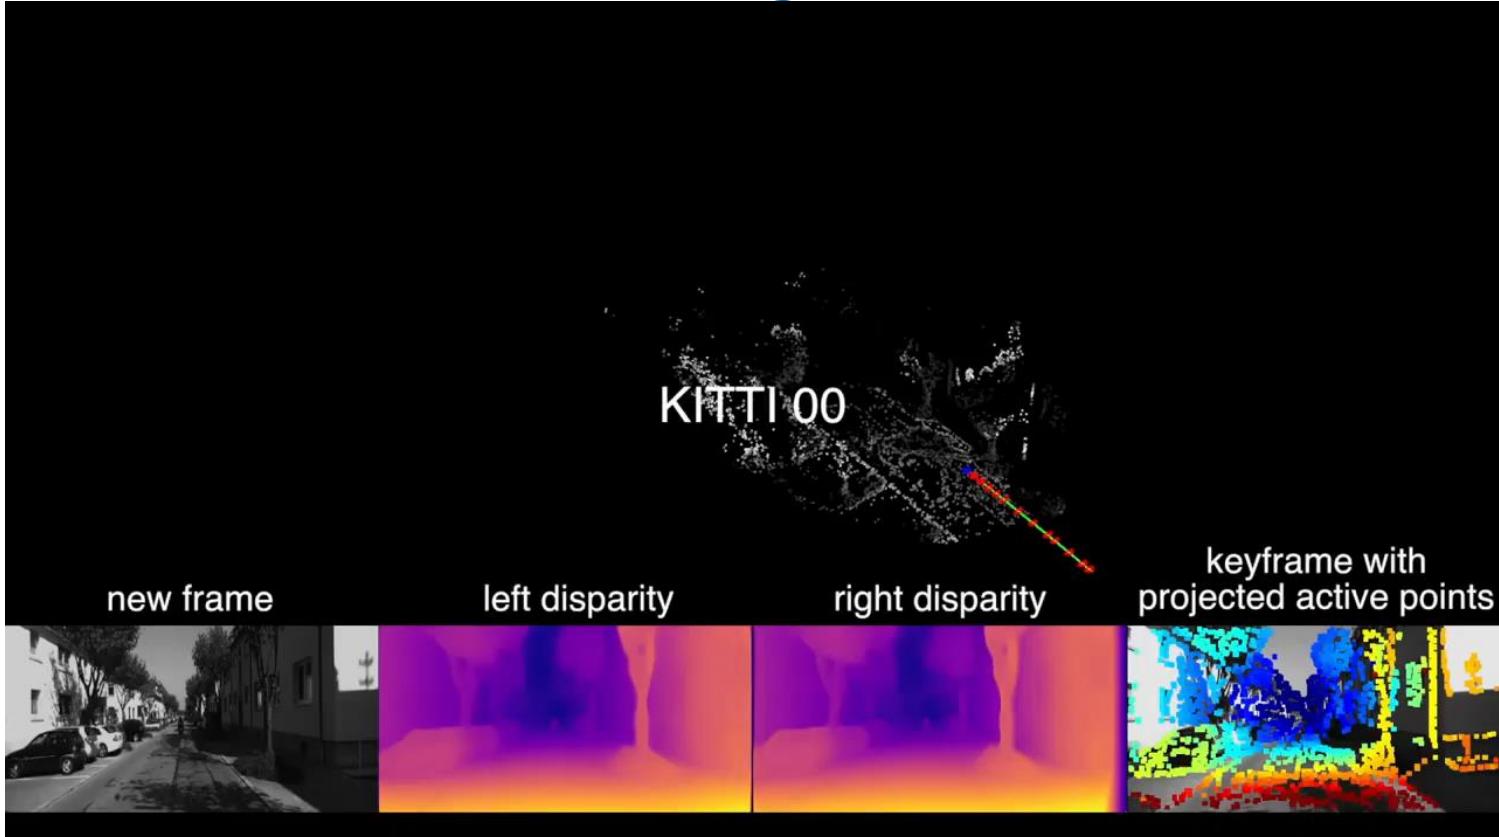
Deep Learning at TUM

- Multiple object tracking with graph neural networks



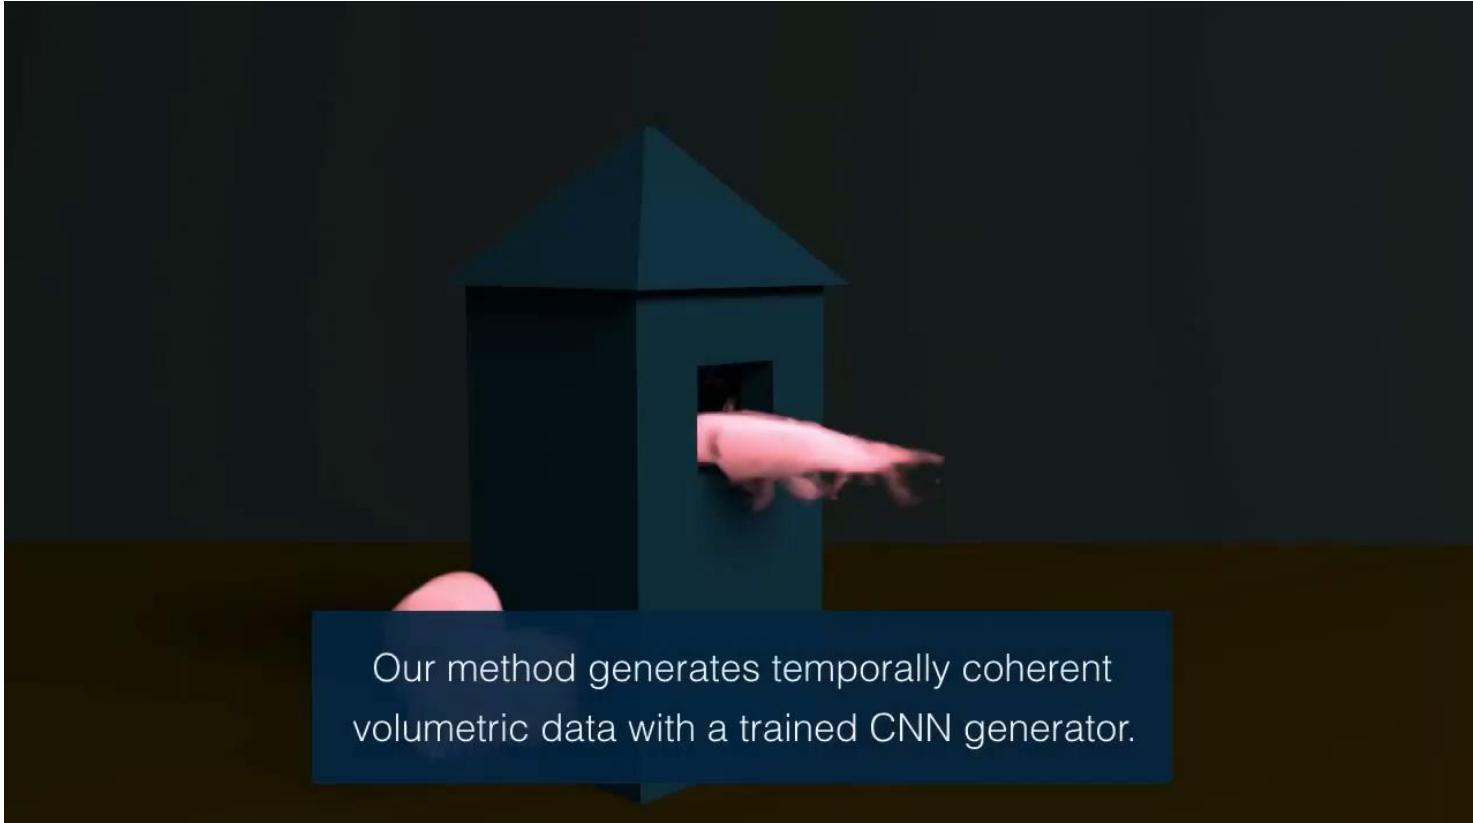
[Brasó and Leal-Taixé, CVPR 2020] Learning a Neural Solver for Multiple Object Tracking.

Deep Learning at TUM



[Yang et al., ECCV'18] Deep Virtual Stereo Odometry
Introduction to Deep Learning

Deep Learning at TUM



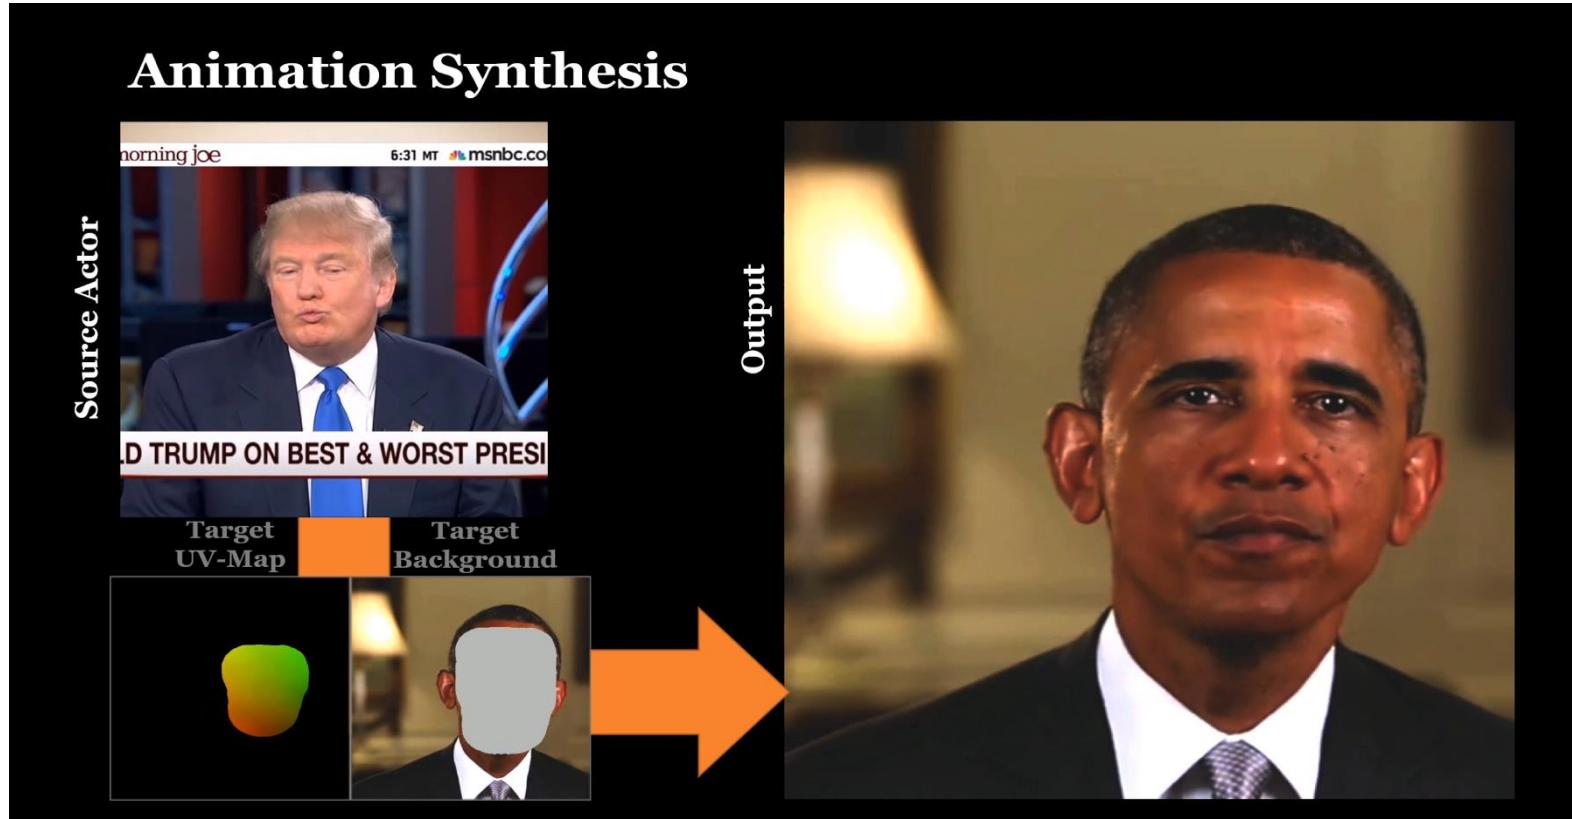
Our method generates temporally coherent
volumetric data with a trained CNN generator.

[Xie et al. Siggraph' 18] tempoGAN

Introduction to Deep Learning

Deep Learning at TUM

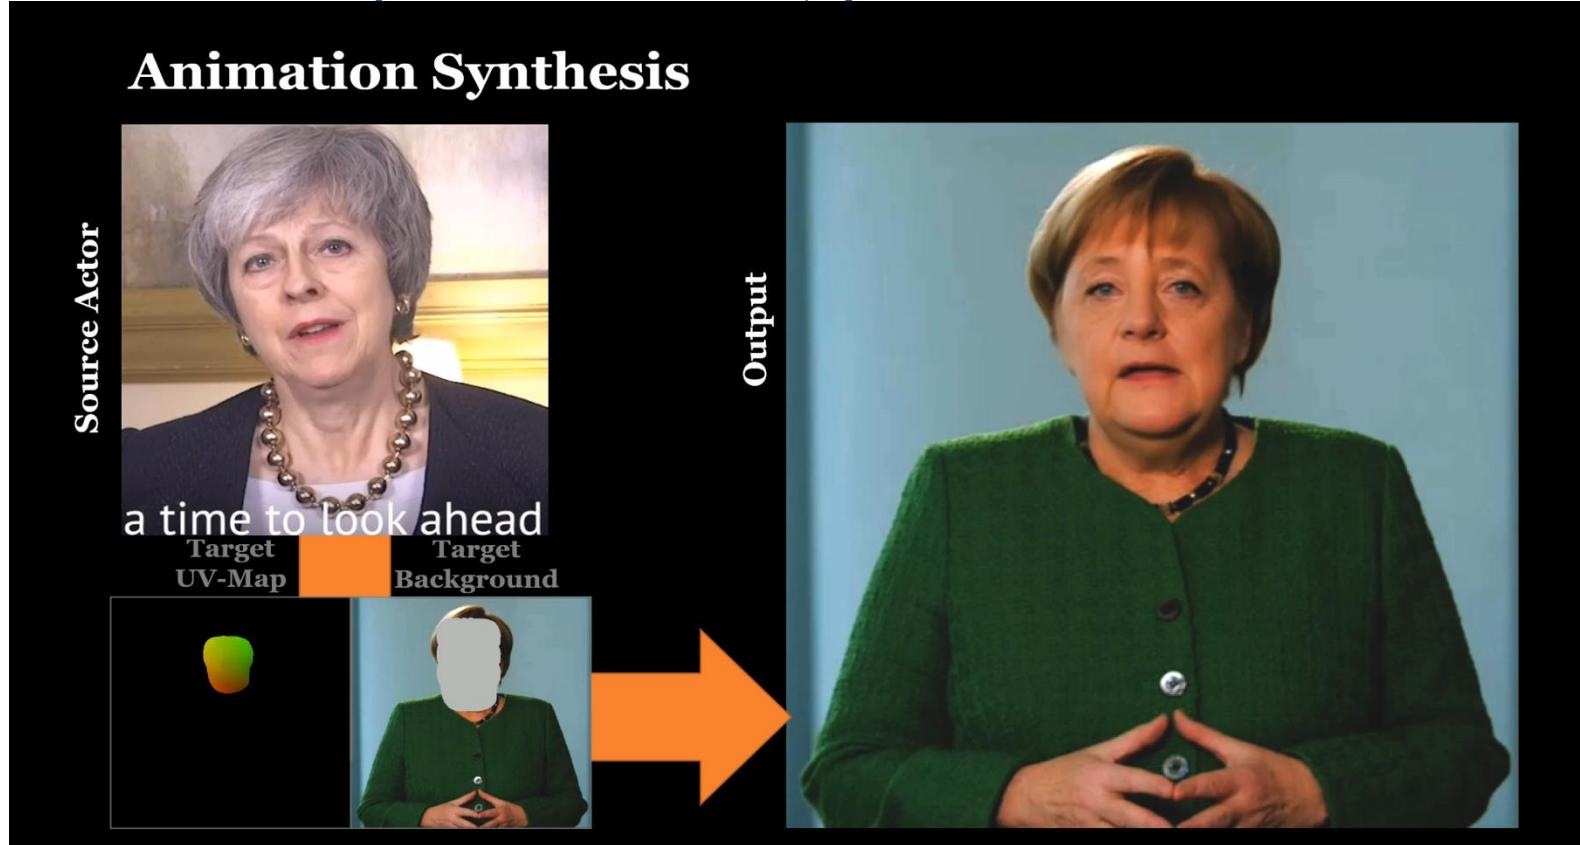
Animation Synthesis



[Thies et al., Siggraph'19]: Neural Textures
Introduction to Deep Learning

Deep Learning at TUM

Animation Synthesis



[Thies et al., Siggraph'19]: Neural Textures

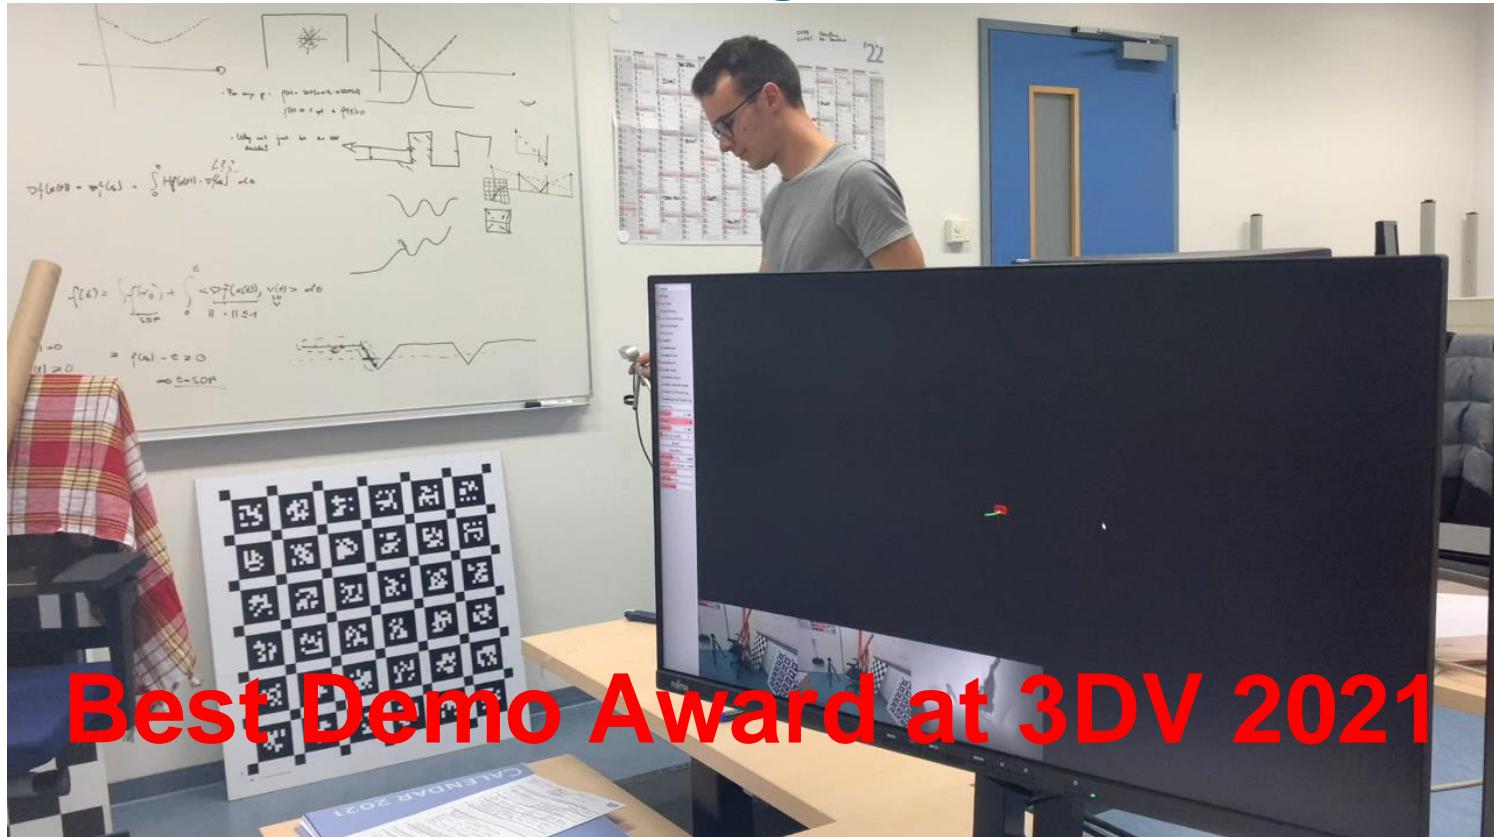
Introduction to Deep Learning

Deep Learning at TUM



Wimbauer et al., "MonoRec: Monocular Dense Reconstruction", CVPR '21
Introduction to Deep Learning

Deep Learning at TUM



Best Demo Award at 3DV 2021

Köstler et al., “TANDEM: Tracking and Dense Mapping”, CoRL ‘21
Introduction to Deep Learning

Deep Learning at TUM

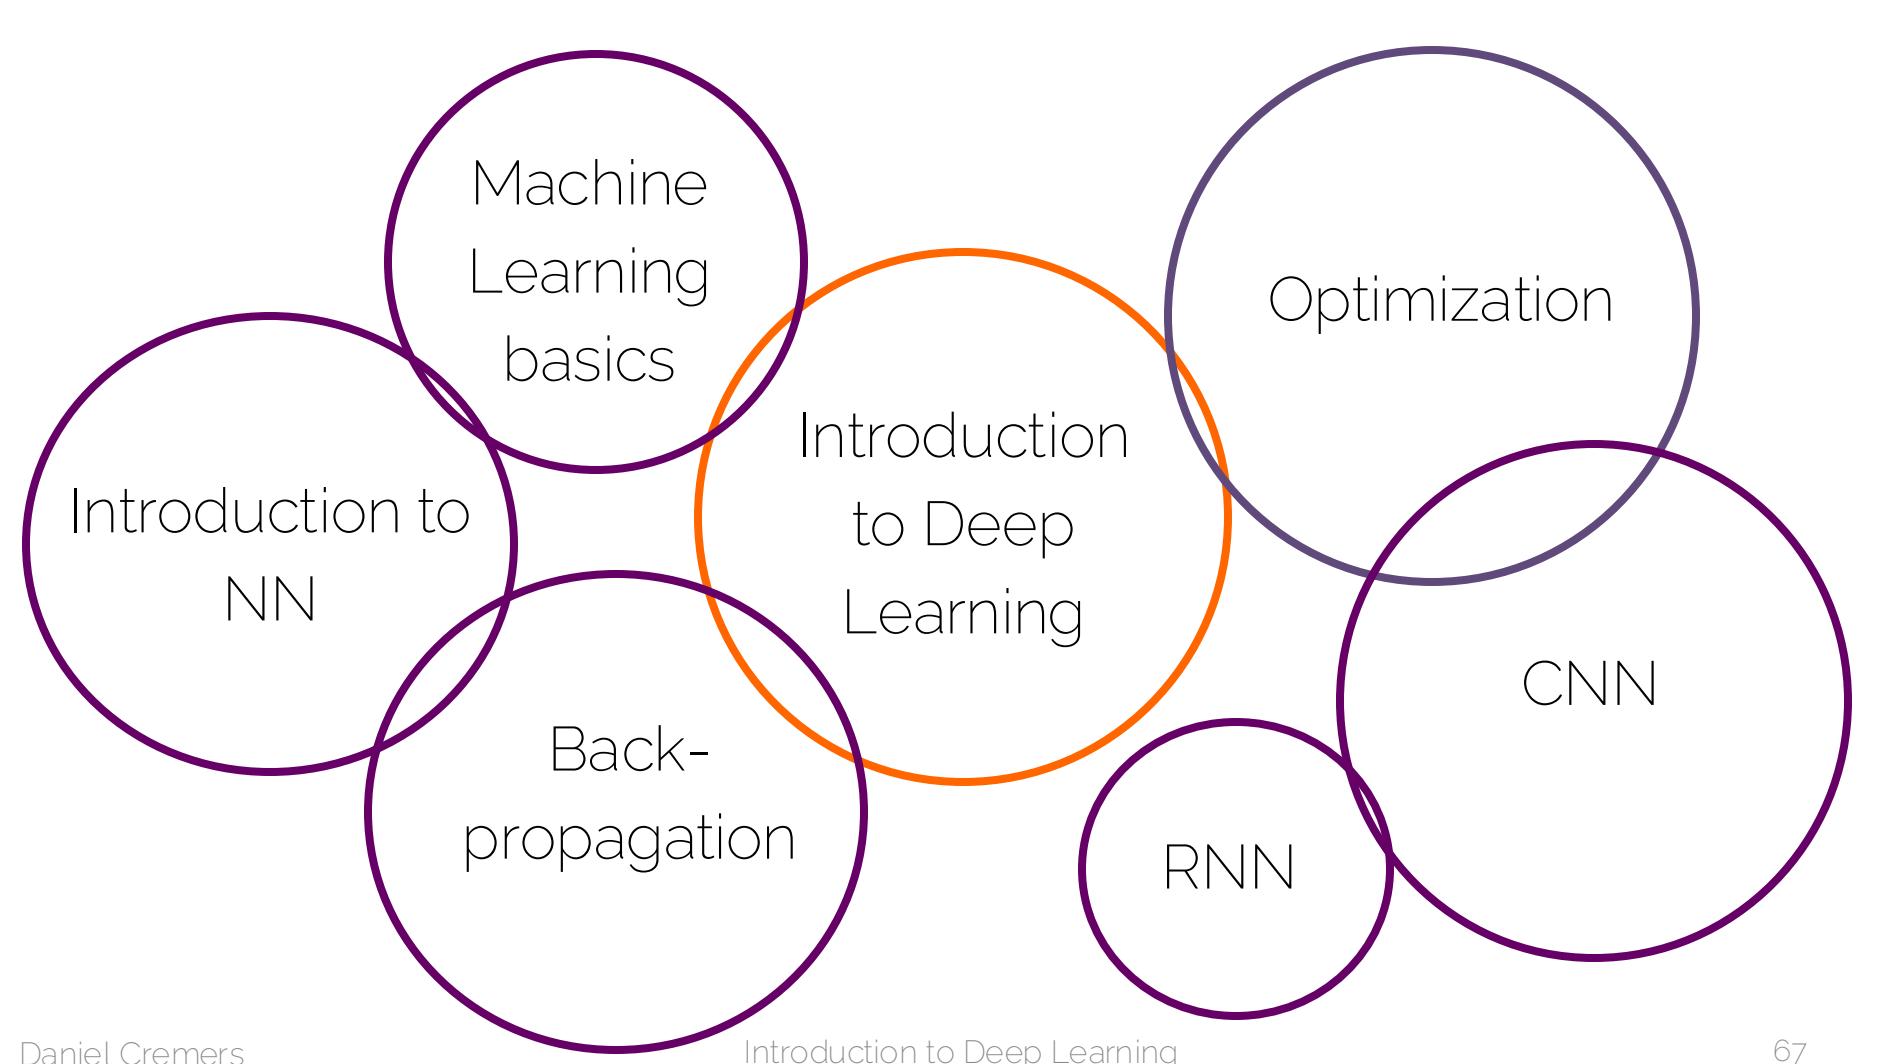


ScanNet Stats:

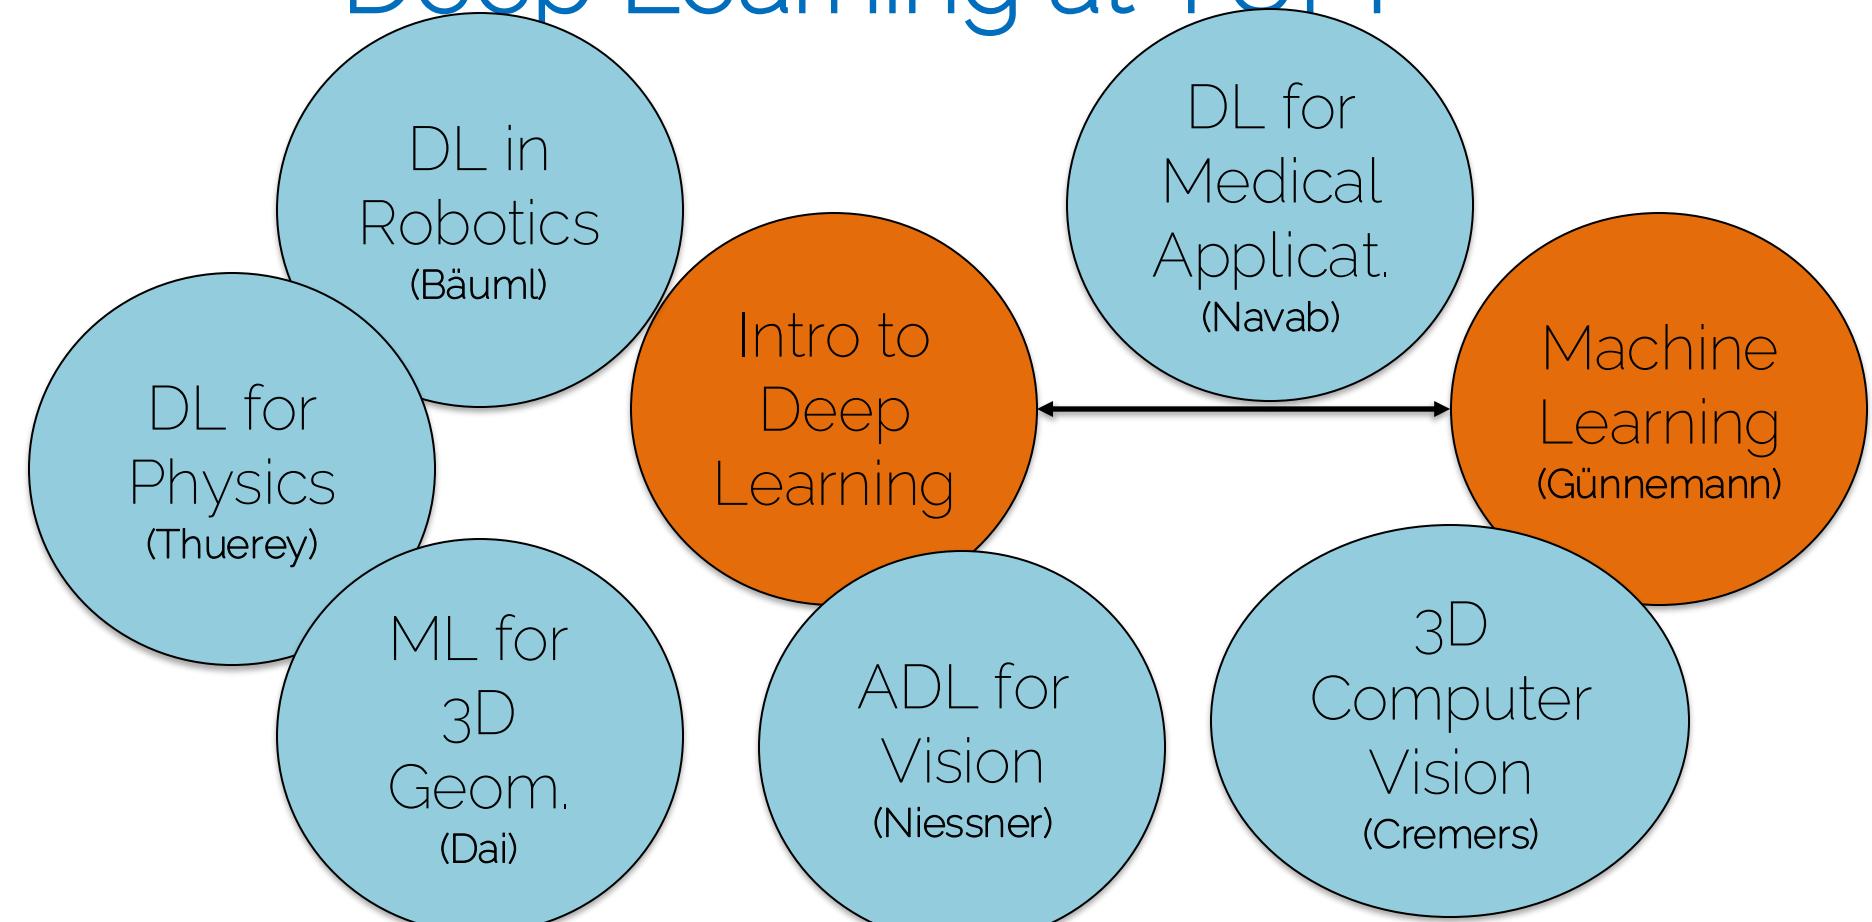
- Kinect-style RGB-D sensors
- 1513 scans of 3D environments
- 2.5 Mio RGB-D frames
- Dense 3D, crowd-source MTurk labels
- Annotations projected to 2D frames

[Dai et al., CVPR'17] ScanNet

Relation to other Lectures at TUM



Deep Learning at TUM



Introduction to Deep Learning

Logistics

The Team

Lecturer



Prof. Dr. Daniel
Cremers

TAs



Shenhan
Qian



Jason
Chui



Yan
Xia

Student Tutors



Julius
Horsch



Oleg
Magnes



Luca
Fehling-Schuh



Subhan
Kamilov



Benjamin
Heltzel



Alexandra
Samoylova



Yunxiang
Lu

About the Lecture

- Theory lectures (every Tuesday at 14:00)
 - In-person, live-streamed
- Tutorials and exercises (every Thursday at 10:00)
 - Tutorial: Online videos posted to Piazza and the webpage
 - Practical exercises
- Guest lecture!

Preliminary Syllabus

Lecture 1: Introduction to the lecture, Deep Learning, Machine Learning.

Lecture 2: Machine Learning Basics, Linear regression, Maximum Likelihood

Lecture 3: Introduction to Neural Networks, Computational Graphs

Lecture 4: Optimization and Backpropagation

Lecture 5: Scaling Optimization to large Data, Stochastic Gradient Descent

Lecture 6: Training Neural Networks I

Lecture 7: Training Neural Networks II

Lecture 8: Training Neural Networks III

Lecture 9: Introduction to CNNs

Lecture 10: CNNs architectures;

Lecture 11: Recurrent Neural Networks (RNNs)

Lecture 12: Advanced Deep Learning architectures

Moodle → Piazza

- Announcements via Piazza - **IMPORTANT!**
 - Sign up online for access: <http://piazza.com/tum.de>
 - Select "Fall 2024" term, search for IN2346
 - Use your @mytum.de email address
 - We will share common information (e.g., regarding exam)
- Forum
 - Ask and discuss questions
 - Tutors will monitor and answer questions
 - You are very welcome to actively participate
 - Please do not post solutions of the exercises
 - You can post private question visible only to the staff



Website

Computer Vision Group
TUM School of Computation, Information and Technology
Technical University of Munich



Home

[Application Form](#)

[TUM AI Lecture Series](#)

[Members](#)

[Research Areas](#)

[Publications](#)

Teaching

[Winter Semester
2024/25](#)

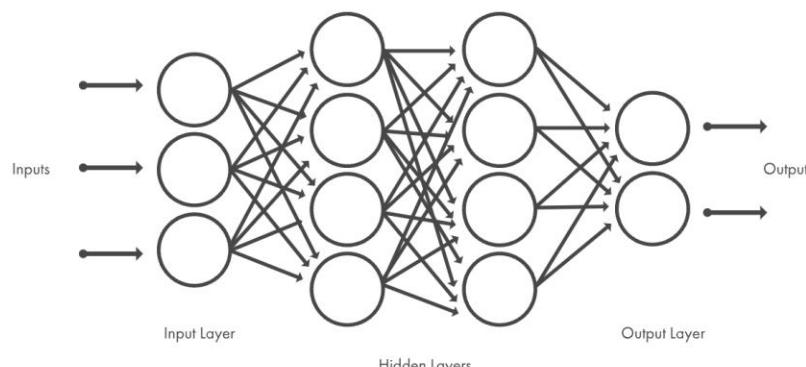
[Computer Vision III:
Detection, Segmentation
and Tracking \(IN2375\)](#)

[Introduction to Deep
Learning \(IN2346\)](#)

[Master Seminar - Beyond
Deep Learning \(5 ECTS\)](#)

[Master Seminar - Recent
Advances in 4D](#)

Introduction to Deep Learning (IN2346)



Informatik IX
Computer Vision Group

Boltzmannstrasse 3
85748 Garching
info@vision.in.tum.de

Follow us on:

[YouTube](#)

[X / Twitter](#)

[Facebook](#)

News

03.07.2024

We have [seven papers](#) accepted
to [ECCV 2024](#). Check our
[publication page](#) for more details.

09.06.2024

[GCPR / VMV 2024](#)

We are organizing [GCPR /](#)
[VMV 2024](#) this fall.

<https://cvg.cit.tum.de/teaching/ws2024/i2dl>

Email

- Email list:

i2dl@vision.in.tum.de

- Do NOT email us personally!
 - Cannot handle so many emails / hence will be ignored
- Email list for organizational questions only!
 - Content questions -> Piazza or Office Hours
 - Or post the question/issue in a private thread on Piazza

(Virtual) Office Hours

- We will have dedicated office hours regarding
 - Theoretical help (e.g., specific lecture questions)
 - Help on exercises
- More info in the first tutorial session
- Zoom links will be posted on Piazza

Exam FAQ

- Final Exam: TBA
- Content: Lecture & exercises
- Important: No retake exam (I2DL is taught every semester)
- Grade Bonus:
 - Solve 9 out of 10 “non-optional” programming exercises
 - Bonus 0.3 on a **passed** final exam
 - Bonus is transferable from previous and future semesters

Other Administrative

- "External" students welcome (LMU, TUM PhD)
 - Fill out registration form and we will add you to the course
 - Will get Certificate / Schein at the end
- Again:
 - Check announcements on Piazza
 - Check content on website:
<https://cvg.cit.tum.de/teaching/ws2024/i2dl>

See you next time 😊

Upcoming Lecture

- Next Lecture: Lecture 2: Machine Learning basics
- Thursday: Tutorial 1 and Exercise 1