

Advanced DL Topics

Attention

Index the values
via a differentiable
operator.


Multiply queries
with keys

Get the values

$$\text{Attention}(Q, K, V) = \text{softmax} \left(\frac{QK^T}{\sqrt{d_k}} \right) V$$

To train them well, divide by $\sqrt{d_k}$, "probably" because for large values of the key's dimension, the dot product grows large in magnitude, pushing the softmax function into regions where it has extremely small gradients.

Transformers




Attention Is All You Need [Vaswani et al. 17]

<https://arxiv.org/pdf/1706.03762.pdf>

Graph Neural Networks

A graph

- Node: a concept
- Edge: a connection between concepts



Deep learning on graphs

- Generalizations of neural networks that can operate on graph-structured domains:
 - Scarselli et al. "The Graph Neural Network Model", IEEE Trans. Neur. Net 2009.
 - Defferrard et al. "Convolutional Neural Networks on Graphs with Fast Localized Spectral Filtering", NeurIPS 2016
 - Kipf&Welling, "Semi-Supervised Classification with Graph Convolutional Networks", ICLR 2017.
 - Gilmer et al. "Neural Message Passing for Quantum Chemistry". ICML 2017
 - Koke&Cremers "HoloNets: Spectral Convolutions do extend to Directed Graphs", ICLR 2024.
- Key challenges:
 - Variable sized inputs (number of nodes and edges)
 - Need **invariance to node permutations**

General Idea 1: Message passing




Figure credit: <https://tkipf.github.io/graph-convolutional-networks/>

General Idea 1: Message passing





Figure credit: <https://tkipf.github.io/graph-convolutional-networks/>


Introduction to Deep Learning

Message Passing Networks


- We can divide the propagation process in two steps: 'node to edge' and 'edge to node' updates.



Initial Graph



'Node to Edge' Update



'Edge to Node' Update

Node embeddings
 Edge embeddings

Battaglia et al. "Relational inductive biases, deep learning, and graph networks". 2018

'Node to edge' updates

- At every message passing step l , first do:

$$h_{(i,j)}^{(l)} = \mathcal{N}_e \left([h_i^{(l-1)}, h_{(i,j)}^{(l-1)}, h_j^{(l-1)}] \right)$$

The diagram shows the equation for 'Node to edge' updates. Three orange arrows point from the text descriptions below to the corresponding terms in the equation. The first arrow points to $h_i^{(l-1)}$, the second to $h_{(i,j)}^{(l-1)}$, and the third to $h_j^{(l-1)}$.

Embedding of node i in the previous message passing step


Embedding of edge (i,j) in the previous message passing step

Embedding of node j in the previous message passing step

'Node to edge' updates

- At every message passing step l , first do:

$$h_{(i,j)}^{(l)} = \mathcal{N}_e \left([h_i^{(l-1)}, h_{(i,j)}^{(l-1)}, h_j^{(l-1)}] \right)$$




'Node to edge' updates

- At every message passing step l , first do:

$$h_{(i,j)}^{(l)} = \underbrace{\mathcal{N}_e}_{\text{Learnable function}} \left([h_i^{(l-1)}, h_{(i,j)}^{(l-1)}, h_j^{(l-1)}] \right)$$

Learnable function (e.g.
MLP) with shared
weights across the
entire graph




'Edge to node' updates

- After a round of edge updates, each edge embedding contains information about its pair of incident nodes
- Then, edge embeddings are used to update nodes:

$$m_i^{(l)} = \Phi \left(\left\{ h_{(i,j)}^{(l)} \right\}_{j \in N_i} \right)$$

Order invariant
operation (e.g.
sum, mean, max)

Neighbors of
node i



'Edge to node' updates


- After a round of edge updates, each edge embedding contains information about its pair of incident nodes
- Then, edge embeddings are used to update nodes:

$$m_i^{(l)} = \Phi \left(\left\{ h_{(i,j)}^{(l)} \right\}_{j \in N_i} \right)$$
$$h_i^{(l)} = \mathcal{N}_v \left([m_i^{(l)}, h_i^{(l-1)}] \right)$$



The aggregation provides each node embedding with contextual information about its neighbors

Learnable function (e.g. MLP) with shared weights across the entire graph



General Idea 2: Spectral Approach



How to extend Convolutions to Graphs?



How to extend Convolutions to Graphs?



Sliding Window Interpretation

How to extend Convolutions to Graphs?

Spectral Approach!

How to extend Convolutions to Graphs?

Spectral Approach!

Convolution

=

Multiplication in Fourier Domain

How to extend Convolutions to Graphs?

Spectral Approach!

Convolution

=

Multiplication in Fourier Domain

$$\text{Conv. Filter} \sim \mathcal{F}^{-1} [g(k) \cdot \mathcal{F}f]$$

How to extend Convolutions to Graphs?

Spectral Approach!

Convolution

=

Multiplication in Fourier Domain



Fourier Transform

=

Projection onto $\{e^{ikx}\}_k$

$$[\mathcal{F}f](k) = \int e^{ikx} f(x) dx$$

How to extend Convolutions to Graphs?

Spectral Approach!

Convolution

=

Multiplication in Fourier Domain



Fourier Transform

=

Projection onto $\{e^{ikx}\}_k$



$\{e^{ikx}\}_k$ are
eigenfunctions of Δ

$$\Delta e^{ikx} = (-k^2) \cdot e^{ikx}$$

How to extend Convolutions to Graphs?

Spectral Approach!

Convolution

=

Multiplication in Fourier Domain



Fourier Transform

=

Projection onto $\{e^{ikx}\}_k$



$\{e^{ikx}\}_k$ are
eigenfunctions of Δ

How to extend Convolutions to Graphs?

Spectral Approach!

Convolution

=

Multiplication in Fourier Domain



Fourier Transform

=

Projection onto $\{e^{ikx}\}_k$



$\{e^{ikx}\}_k$ are
eigenfunctions of Δ

Take $\{\phi_k\}_k$
eigenvectors of Δ

How to extend Convolutions to Graphs?

Spectral Approach!

Convolution

=

Multiplication in Fourier Domain



Fourier Transform

=

Projection onto $\{e^{ikx}\}_k$



$\{e^{ikx}\}_k$ are
eigenfunctions of Δ

Graph Fourier Transform

:=


Projection onto $\{\phi_k\}_k$



Take $\{\phi_k\}_k$
eigenvectors of Δ


How to extend Convolutions to Graphs?

Spectral Approach!



How to extend Convolutions to Graphs?

Spectral Approach!



General Idea 3: Holomorphic Functional Calculus

- Key idea: Extend spectral convolutions to graphs by making use of
 - complex analysis
 - holomorphic functions & the Cauchy integral formula
 - tools from spectral theory

[Koke & Cremers "HoloNets: Spectral Convolutions do extend to Directed Graphs", ICLR 2024]

General Idea 3: Holomorphic Functional Calculus

Table 1: Results on real-world directed heterophilic datasets. OOM indicates out of memory.

| Homophily | Squirrel 0.223 | Chameleon 0.235 | Arxiv-year 0.221 | Snap-patents 0.218 | Roman-Empire 0.05 |
|--------------------|------------------------------------|------------------------------------|------------------------------------|------------------------------------|------------------------------------|
| MLP | 28.77 ± 1.56 | 46.21 ± 2.99 | 36.70 ± 0.21 | 31.34 ± 0.05 | 64.94 ± 0.62 |
| GCN | 53.43 ± 2.01 | 64.82 ± 2.24 | 46.02 ± 0.26 | 51.02 ± 0.06 | 73.69 ± 0.74 |
| H ₂ GCN | 37.90 ± 2.02 | 59.39 ± 1.98 | 49.09 ± 0.10 | OOM | 60.11 ± 0.52 |
| GPR-GNN | 54.35 ± 0.87 | 62.85 ± 2.90 | 45.07 ± 0.21 | 40.19 ± 0.03 | 64.85 ± 0.27 |
| LINKX | 61.81 ± 1.80 | 68.42 ± 1.38 | 56.00 ± 0.17 | 61.95 ± 0.12 | 37.55 ± 0.36 |
| FSGNN | 74.10 ± 1.89 | 78.27 ± 1.28 | 50.47 ± 0.21 | 65.07 ± 0.03 | 79.92 ± 0.56 |
| ACM-GCN | 67.40 ± 2.21 | 74.76 ± 2.20 | 47.37 ± 0.59 | 55.14 ± 0.16 | 69.66 ± 0.62 |
| GloGNN | 57.88 ± 1.76 | 71.21 ± 1.84 | 54.79 ± 0.25 | 62.09 ± 0.27 | 59.63 ± 0.69 |
| Grad. Gating | 64.26 ± 2.38 | 71.40 ± 2.38 | 63.30 ± 1.84 | 69.50 ± 0.39 | 82.16 ± 0.78 |
| DiGCN | 37.74 ± 1.54 | 52.24 ± 3.65 | OOM | OOM | 52.71 ± 0.32 |
| MagNet | 39.01 ± 1.93 | 58.22 ± 2.87 | 60.29 ± 0.27 | OOM | 88.07 ± 0.27 |
| DirGNN | 75.13 ± 1.95 | 79.74 ± 1.40 | 63.97 ± 0.30 | 73.95 ± 0.05 | 91.3 ± 0.46 |
| FaberNet | 76.71 ± 1.92 | 80.33 ± 1.19 | 64.62 ± 1.01 | 75.10 ± 0.03 | 92.24 ± 0.43 |

[Koke & Cremers "HoloNets: Spectral Convolutions do extend to Directed Graphs", ICLR 2024]

GNN Applications


- Node or edge classification
 - identifying anomalies such as spam, fraud
 - Relationship discovery for social networks, search networks



Image size:
881 × 657

[Find other sizes of this image](#)
[All sizes](#) · [Small](#) · [Medium](#) · [Large](#)

[Visually similar images](#)




<https://gm-neurips-2020.github.io/master-deck.pdf>

Introduction to Deep Learning

GNN Applications


- Scene graph generation



[Xu et al. '17] Scene Graph Generation by Iterative Message Passing


GNN Applications

- 3D Mesh Classification



GNN Applications


- 3D mesh generation



[Dai and Niessner, "Scan2Mesh: From Unstructured Range Scans to 3D Meshes", CVPR 2019]

GNN Applications

- Modeling epidemiology
 - Spatio-temporal graph

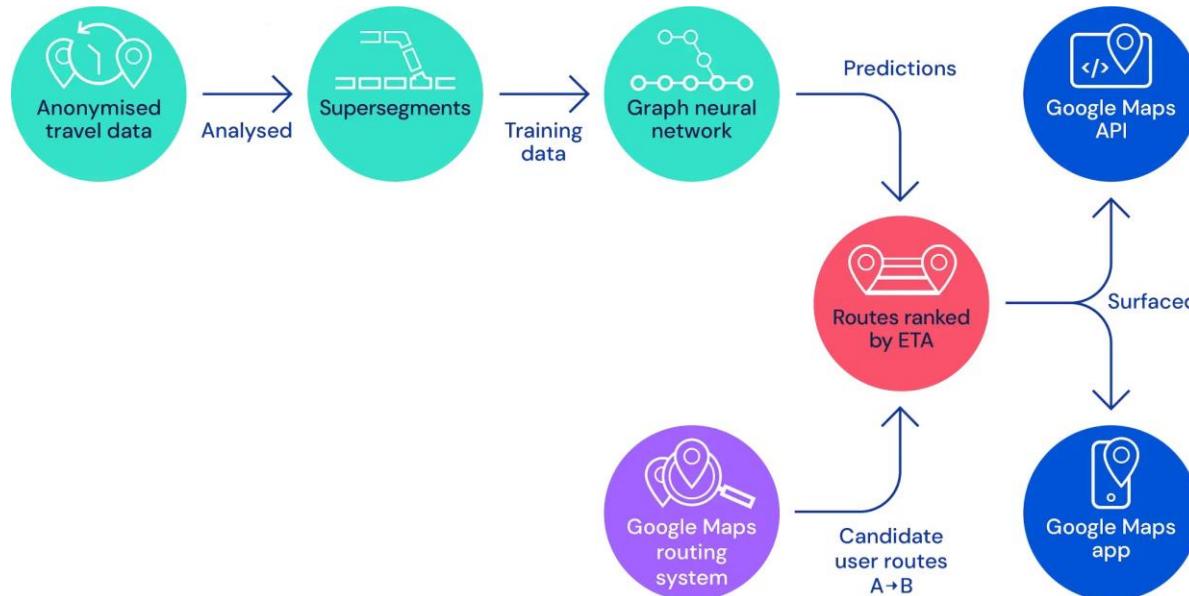


<https://gm-neurips-2020.github.io/master-deck.pdf>

Introduction to Deep Learning

GNN Applications

- Traffic forecasting




<https://www.deeplearning.ai/ai-notes/traffic-prediction-with-graph-neural-networks/>

Generative Models

Generative Models

- Given training data, how to generate new samples from the same distribution

Real Images



Generated Images



Source: <https://openai.com/blog/generative-models/>

Generative Models





Figure copyright and adapted from Ian Goodfellow, Tutorial on Generative Adversarial Networks, 2017

Autoencoders

Autoencoders


- Can be used as a basic generative models
- Unsupervised approach for learning a lower-dimensional feature representation from unlabeled training data

Autoencoders




- From an input image to a feature representation (bottleneck layer)
- Encoder: a CNN in our case


Autoencoder: training



Autoencoder: training




Autoencoder: training




- No labels required
- We can use unlabeled data to first get its structure

Autoencoder



Decoder as Generative Model

Test time:
-> reconstruction from
'random' vector




Why using autoencoders?

- Use 1: pre-training, as mentioned before
 - Image → same image reconstructed
 - Use the encoder as "feature extractor"
- Use 2: Use them to get pixel-wise predictions
 - Image → semantic segmentation
 - Low-resolution image → High-resolution image
 - Image → Depth map

Variational Autoencoders


Autoencoders

- Encode the input into a representation (bottleneck) and reconstruct it with the decoder



Autoencoders


- Encode the input into a representation (bottleneck) and reconstruct it with the decoder



Latent space learned
by autoencoder on MNIST


Source: <https://bit.ly/37ctFMS>

Variational Autoencoder




Variational Autoencoder

Goal: Sample from the latent distribution to generate new outputs!




Variational Autoencoder

- Latent space is now a distribution
- Specifically it is a Gaussian




Variational Autoencoder

- Latent space is now a distribution
- Specifically it is a Gaussian




Variational Autoencoder

- Training: loss makes sure the latent space is close to a Gaussian and the reconstructed output is close to the input



Variational Autoencoder

- Test: Sample from the latent space



Autoencoder vs VAE



Autoencoder



Variational Autoencoder



Ground Truth

Source: <https://github.com/kvfrans/variational-autoencoder>

Generating data

Degree of smile




Head pose

Autoencoder Overview

- Autoencoders (AE)
 - Reconstruct input
 - Unsupervised learning
- Variational Autoencoders (VAE)
 - Probability distribution in latent space (e.g., Gaussian)
 - Interpretable latent space (head pose, smile)
 - Sample from model to generate output


Generative Adversarial Networks (GANs)

Generative Adversarial Networks (GANs)




Source: <https://github.com/hindupuravinash/the-gan-zoo>

Autoencoder




Decoder as Generative Model

Test time:
-> reconstruction from
'random' vector




Decoder as Generative Model

"Test time":
-> reconstruction from
'random' vector




Reconstruction Loss
Often L2, i.e., sum of squared dist.
-> L2 distributes error equally
-> mean is opt.
-> res. is blurry




Instead of L2, can we
"learn" a loss function?

Generative Adversarial Networks (GANs)




[Goodfellow et al., NIPS'14] Generative Adversarial Networks (slide from McGuinness)

Generative Adversarial Networks (GANs)



[Goodfellow et al., NIPS'14] Generative Adversarial Networks (slide from McGuinness)

Generative Adversarial Networks (GANs)



(Goodfellow 2016)

[Goodfellow, NIPS'16] Tutorial: Generative Adversarial Networks
Introduction to Deep Learning

GANs: Loss Functions

- Discriminator loss

$$J^{(D)} = -\frac{1}{2} \mathbb{E}_{\mathbf{x} \sim p_{data}} \log D(\mathbf{x}) - \frac{1}{2} \mathbb{E}_{\mathbf{z}} \log (1 - D(G(\mathbf{z})))$$

binary cross entropy

- Generator loss


$$J^{(G)} = -J^{(D)}$$

- Minimax Game:

- G minimizes probability that D is correct
 - Equilibrium is saddle point of discriminator loss
 - D provides supervision (i.e., gradients) for G

GAN Applications

BigGAN: HD Image Generation



[Brock et al., ICLR'18] BigGAN : Large Scale GAN Training for High Fidelity Natural Image Synthesis


StyleGAN: Face Image Generation



[Karras et al., '18] StyleGAN : A Style-Based Generator Architecture for Generative Adversarial Networks



[Karras et al., '19] StyleGAN2 : Analyzing and Improving the Image Quality of StyleGAN

Cycle GAN: Unpaired Image-to-Image Translation






[Zhu et al., ICCV'17] Cycle GAN : Unpaired Image-to-Image Translation using Cycle-Consistent Adversarial Networks
Daniel Cremers

SPADE: GAN-Based Image Editing



Texturify: 3D Texture Generation



[Siddiqui et al., ECCV'22]

74

Diffusion

Diffusion – Search Interest

Interest over time 




Source: Google Trends

Diffusion Models

- Class of generative models
- Achieved state-of-the-art image generation (DALLE-2, Imagen, StableDiffusion)
- What is diffusion?

Diffusion Process

- Gradually add noise to input image x_0 in a series of T time steps
- Neural network trained to recover original data



[Ho et al. '20] Denoising Diffusion Probabilistic Models


Introduction to Deep Learning

Forward Diffusion

- Markov chain of T steps
 - Each step depends only on previous
 - Adds noise to x_0 sampled from real distribution $q(x)$

$$q(x_t | x_{t-1}) = \mathcal{N}(x_t; \mu_t = \sqrt{1 - \beta_t} x_{t-1}, \Sigma_t = \beta_t \mathbf{I})$$

mean variance identity matrix



[Ho et al. '20] Denoising Diffusion Probabilistic Models
Introduction to Deep Learning

Forward Diffusion

- Go from x_0 to x_T :

$$q(x_{1:T}|x_0) = \prod_{t=1}^T q(x_t|x_{t-1})$$

- Efficiency?

Reparameterization

- Define $\alpha_t = 1 - \beta_t$, $\bar{\alpha}_t = \prod_{s=0}^t \alpha_s$, $\epsilon_0, \dots, \epsilon_{t-1} \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$

$$\begin{aligned}x_t &= \sqrt{1 - \beta_t} x_{t-1} + \sqrt{\beta_t} \epsilon_{t-1} \\&= \sqrt{\alpha_t} x_{t-1} + \sqrt{1 - \alpha_t} \epsilon_{t-1} \\&= \sqrt{\alpha_t \alpha_{t-1}} x_{t-2} + \sqrt{1 - \alpha_t \alpha_{t-1}} \epsilon_{t-2} \\&= \sqrt{\bar{\alpha}_t} x_0 + \sqrt{1 - \bar{\alpha}_t} \epsilon_0\end{aligned}$$

$$x_t \sim q(x_t | x_0) = \mathcal{N}(x_t; \sqrt{\bar{\alpha}_t} x_0, (1 - \bar{\alpha}_t) \mathbf{I})$$

Reverse Diffusion


- $x_{T \rightarrow \infty}$ becomes a Gaussian distribution
- Reverse distribution $q(x_{t-1}|x_t)$
 - Sample $x_T \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$ and run reverse process
 - Generates a novel data point from original distribution
- How to model reverse process?

Approximate Reverse Process

- Approximate $q(x_{t-1}|x_t)$ with parameterized model p_θ (e.g., deep network)

$$p_\theta(x_{t-1}|x_t) = \mathcal{N}(x_{t-1}; \mu_\theta(x_t, t), \Sigma_\theta(x_t, t))$$

$$p_\theta(x_{0:T}) = p_\theta(x_T) \prod_{t=1}^T p_\theta(x_{t-1}|x_t)$$



[Ho et al. '20] Denoising Diffusion Probabilistic Models
Introduction to Deep Learning

Training a Diffusion Model

- Optimize negative log-likelihood of training data

$$\begin{aligned} L_{VLB} &= \mathbb{E}_q \left[D_{KL}(q(x_T|x_0) || p_\theta(x_T)) \right]_{L_T} \\ &+ \sum_{t=2}^T \underbrace{D_{KL}(q(x_{t-1}|x_t, x_0) || p_\theta(x_{t-1}|x_t))}_{L_{t-1}} - \underbrace{\log p_\theta(x_0|x_1)}_{L_0} \end{aligned}$$

- Nice derivations: <https://lilianweng.github.io/posts/2021-07-11-diffusion-models>

Training a Diffusion Model

- $L_{t-1} = D_{KL}(q(x_{t-1}|x_t, x_0) || p_\theta(x_{t-1}|x_t))$
- Comparing two Gaussian distributions
- $L_{t-1} \propto \|\tilde{\mu}_t(x_t, x_0) - \mu_\theta(x_t, t)\|^2$
- Predicts diffusion posterior mean

Diffusion Model Architecture

- Input and output dimensions must match
- Highly flexible to architecture design
- Commonly implemented with U-Net architectures

Applications for Diffusion Models


- Text-to-image



Oil Painting



Digital Illustration




Hyperrealistic



Cartoon

Applications for Diffusion Models


- Image inpainting & outpainting



<https://github.com/lkwq007/stablediffusion-infinity>
Introduction to Deep Learning

Applications for Diffusion Models

- Text-to-3D Neural Radiance Fields



<https://dreamfusion3d.github.io/>

Realistic 3D Human Motion Generation with Anisotropic Diffusion

Cecilia Curreli^{1,2}

Zhenzhang Ye¹

Dominik Muhle^{1,2}


Riccardo Marin^{1,2}

Abhishek Saroha^{1,2}

Daniel Cremers^{1,2}

¹Technical University of Munich


²Munich Center for Machine Learning



3D Human Motion Prediction

Input: body joint positions up to time T

$$\text{Past} \in \mathbb{R}^{P \times J \times 3}$$




Output: body joint positions from time T

$$\text{Future} \in \mathbb{R}^{F \times J \times 3}$$

3D Human Motion Prediction

Input: body joint positions up to time T

$$\text{Past} \in \mathbb{R}^{P \times J \times 3}$$




Output: body joint positions from time T

$$\text{Future} \in \mathbb{R}^{F \times J \times 3}$$

3D Human Motion Prediction

Input: body joint positions up to time T


$$\text{Past} \in \mathbb{R}^{P \times J \times 3}$$



Output: body joint positions from time T

$$\text{Future} \in \mathbb{R}^{F \times J \times 3}$$




3D Human Motion Prediction




Anisotropies in Latent Space

Principal Component Analysis of latent space:

• SkeletonDiffusion
★ GT




Recap: Isotropic Diffusion



$$\Sigma_t = (1 - \alpha_t)\mathbb{I}$$


$$x_t = \sqrt{\alpha_t}x_{t-1} + (1 - \alpha_t)\epsilon \quad \epsilon \sim \mathcal{N}(\mathbf{0}, \mathbb{I})$$

Isotropic vs. Anisotropic Diffusion




$$\Sigma_t = (1 - \alpha_t)\mathbb{I}$$

Isotropic vs. Anisotropic Diffusion



Ours:

$$\Sigma_t = (1 - \alpha_t)\mathbb{I}$$




$$\Sigma_t = (1 - \alpha_t)\gamma_t \Sigma_N + (1 - \alpha_t)(1 - \gamma_t)\mathbb{I}$$

Isotropic vs. Anisotropic Diffusion

Ours:

$$\Sigma_N = \frac{\mathbf{A} - \lambda_{\min}(\mathbf{A})\mathbb{I}}{\lambda_{\max}(\mathbf{A}) - \lambda_{\min}(\mathbf{A})}$$

A: adjacency matrix of the skeleton graph



$$\Sigma_t = (1 - \alpha_t)\gamma_t \Sigma_N + (1 - \gamma_t)(1 - \alpha_t)\mathbb{I}$$

Comparison to Baselines

GT



DiverseSampling



CoMusion



Ours



Diverse yet realistic

Comparison to Baselines

GT



DiverseSampling



CoMusion



Ours



Future coherent with past motion

Reinforcement Learning

Learning Paradigms in ML

Supervised Learning

E.g., classification,
regression

Labeled data

Find mapping from
input to label

Unsupervised Learning

E.g., clustering,
anomaly detection

Unlabeled data

Find structure in data


Reinforcement Learning

Sequential data

Learning by
interaction with
the environment


In a Nutshell

- RL-agent is trained using the “carrot and stick” approach
- Good behavior is encouraged by rewards
- Bad behavior is discouraged by punishment



Source: quora.com

Agent and Environment



Characteristics of RL

- Sequential, non i.i.d. data (time matters)
- Actions have an effect on the environment
-> Change future input
- No supervisor, target is approximated by the reward signal

History and State

- The agent makes decisions based on the **history h** of observations, actions and rewards up to time-step t

$$h_t = o_1, a_1, r_1, \dots, a_{t-1}, r_{t-1}, o_t$$

- The **state s** contains all the necessary information from $h \rightarrow s$ is a function of h

$$s_t = f(h_t)$$

Markov Assumption


- Problem: History grows linearly over time
- Solution: **Markov Assumption**
- A state S_t is Markov if and only if:

$$\mathbb{P}[s_{t+1}|s_t] = \mathbb{P}[s_{t+1}|s_1, \dots s_t]$$

- “The future is independent of the past given the present”

Agent and Environment

- Reward and next state are functions of current observation o_t and action a_t only



Mathematical Formulation

- The RL problem is a Markov Decision Process (MDP) defined by: $(\mathcal{S}, \mathcal{A}, \mathcal{R}, \mathbb{P}, \gamma)$

\mathcal{S} : Set of possible states

\mathcal{A} : Set of possible actions

\mathcal{R} : Distribution of reward given (state, action) pair


\mathbb{P} : Transition probability of a (state, action) pair

γ : Discount factor (discounts future rewards)

Components of an RL Agent

- Policy π : Behavior of the agent
-> Mapping from state to action: $a = \pi(s)$
- Value-, Q-Function: How good is a state or (state, action) pair
-> Expected future reward

Taxonomy of RL Algorithms



RL Milestones: Playing Atari



- Mnih et al. 2013, first appearance of DQN
- Successfully learned to play different Atari games like Pong, Breakout, Space Invaders, Seaquest and Beam Rider

[Mnih et al., NIPS'13] Playing Atari with Deep Reinforcement Learning

RL Milestones: AlphaZero (StarCraft)


- Model: Transformer network with a LSTM core
- Trained on 200 years of StarCraft play for 14 days
- 16 Google v3 TPUs
- December 2018:
Beats MaNa, a professional StarCraft player (world rank 13)



I2DL Summary



Machine Learning Basics

- Unsupervised vs Supervised Learning




- Linear vs logistic regression

- Data splitting



Intro to Neural Networks

- Backpropagation




Chain Rule:

$$\frac{\partial f}{\partial x} = \frac{\partial f}{\partial d} \cdot \frac{\partial d}{\partial x}$$

$$\frac{\partial f}{\partial x}$$


- Activation functions



- Loss functions
 - Comparison & effects


Training Neural Networks

- Gradient Descent/ SGD




Source: <http://ruder.io/optimizing-gradient-descent/>,
<https://srdas.github.io/DLBook/ImprovingModelGeneralization.html>, <http://cs231n.github.io/neural-networks-3/>

- Regularization




- Parameter & interpretation




Typology of Neural Networks


- CNNs




- RNNs



- Autoencoder




- GANs



Other DL Courses

Deep Learning at TUM



Next Dates and Exam

- Guest Lecture by Ben Poole!
 - Monday July 17th at 7pm (CEST)
 - Join via Live Stream:
<https://www.youtube.com/watch?v=xk-TibnYEDA>
- Exam
 - There will NOT be a retake exam
 - Neither cheat sheet nor calculator during the exam

Good Luck
in the Exam ☺

References for Further Reading

- <https://towardsdatascience.com/intuitively-understanding-variational-autoencoders-1bfe67eb5daf>
- <https://phillipi.github.io/pix2pix/>
- http://cs231n.stanford.edu/slides/2017/cs231n_2017_lecure13.pdf